

Docker

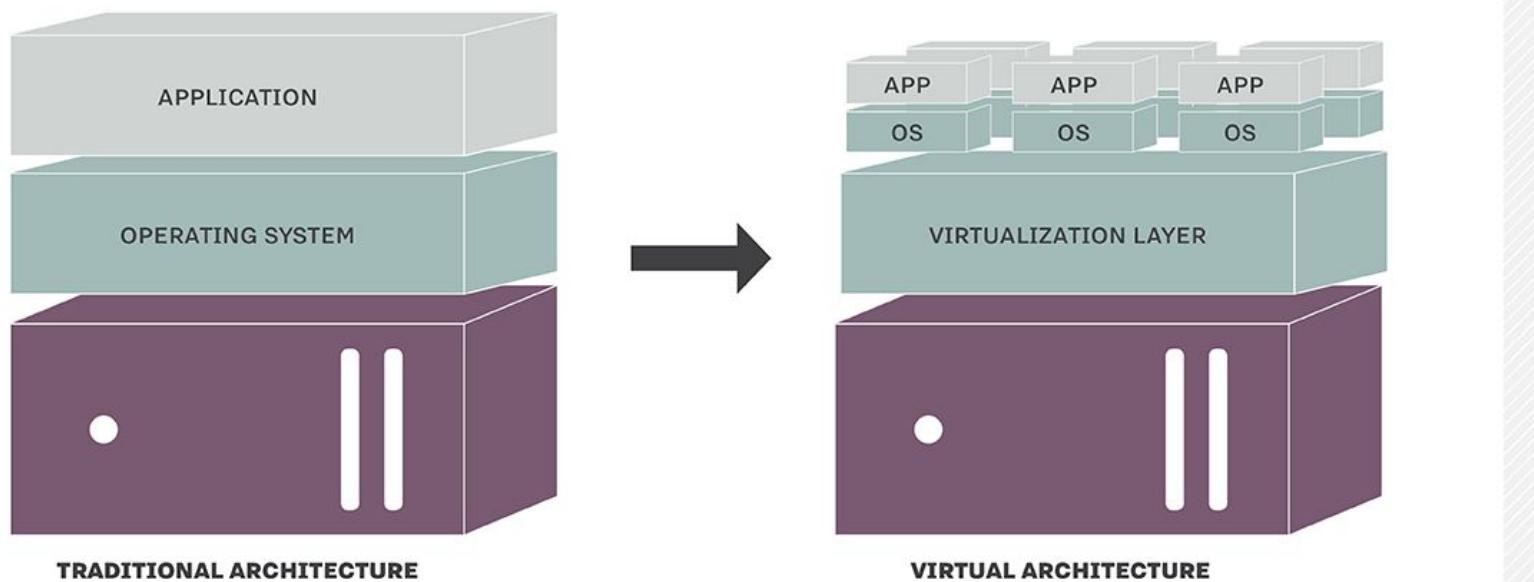
Anurak Theanpurmpul

Agenda

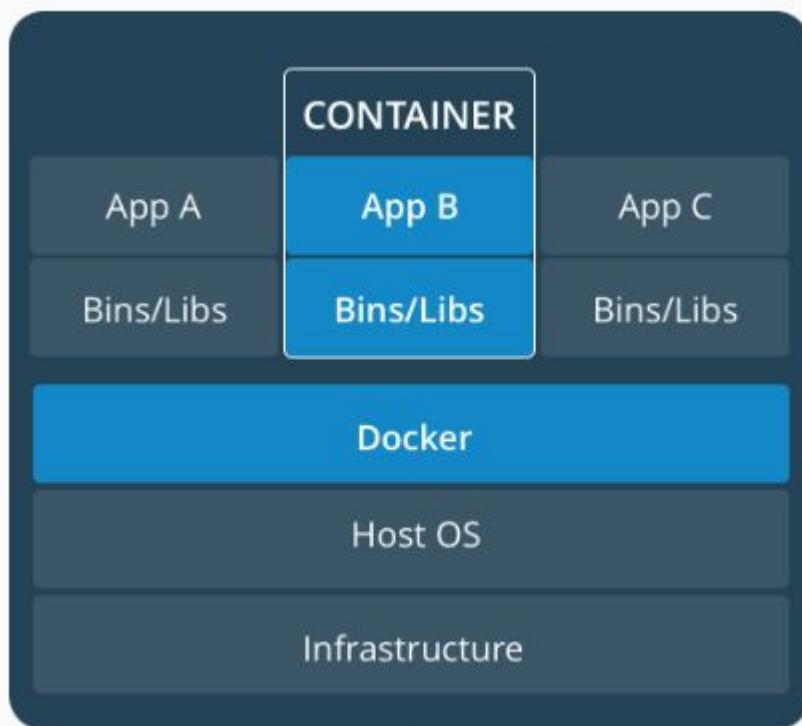
- Docker Overview
- Docker Registry
- Docker Images
- Docker Volume
- Docker Network
- Docker Build (Dockerfile)
- Multi-state Build
- Docker Compose
- Docker Swarm
- Docker Stack

Docker Overview

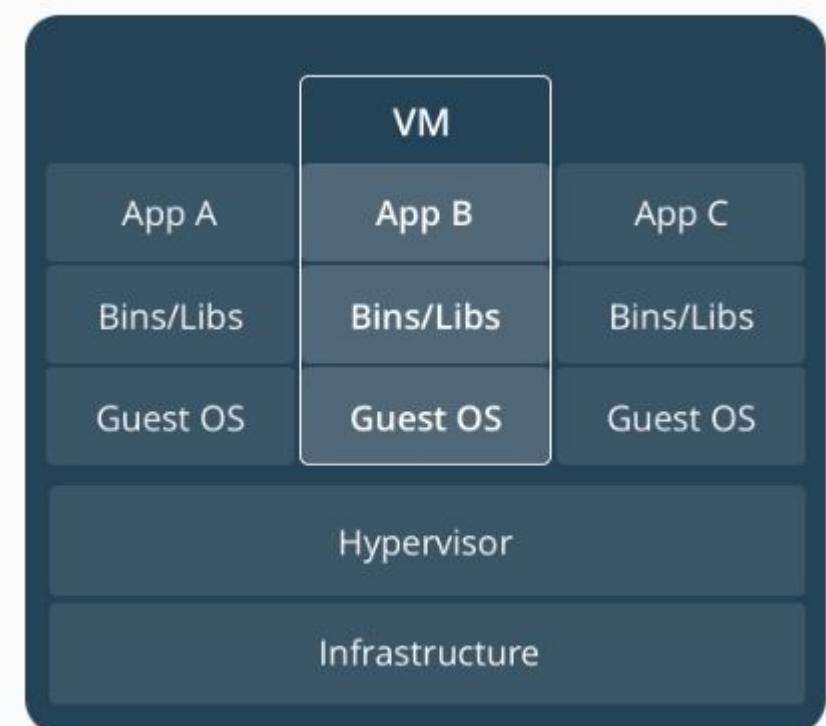
Traditional and Virtual Architecture



Containers vs Virtual Machines



CONTAINERS



VIRTUAL MACHINES

- LXC (LinuX Containers)

Docker

Docker is an open platform for developers and sysadmins to build, ship, and run distributed applications.

<http://www.docker.com/>

Docker

- Docker was born in 2013 by dotCloud, Inc.



dotCloud, Inc. is now Docker, Inc.

Docker, Inc. will be devoting the vast majority of its resources towards growing Docker and the Docker ecosystem [Read More](#)

GIGAOM **InfoQ**

Forbes **InformationWeek**

**Get started with
Docker today.**



Docker Toolbox

- Windows 7,8 or 10 Home (uses Hyper-V)

[Get Docker Toolbox for
Windows](#)

Docker Toolbox is for older Mac and Windows systems that do not meet the requirements of Docker for Mac and Docker for Windows and [Microsoft Hyper-V](#).

Docker Desktop

- Docker on Windows 10 Pro/Ent (uses Hyper-V and check for latest updates)

Get Docker

Docker is an open platform for developing, shipping, and running applications. Docker enables you to separate your applications from your infrastructure so you can deliver software quickly. With Docker, you can manage your infrastructure in the same ways you manage your applications. By taking advantage of Docker's methodologies for shipping, testing, and deploying code quickly, you can significantly reduce the delay between writing code and running it in production.

You can download and install Docker on multiple platforms. Refer to the following section and choose the best installation path for you.



Docker Desktop for Mac

A native application using the macOS sandbox security model which delivers all Docker tools to your Mac.



Docker Desktop for Windows

A native Windows application which delivers all Docker tools to your Windows computer.



Docker for Linux

Install Docker on a computer which already has a Linux distribution installed.

Docker Desktop WSL 2 backend

- No more Hyper-V, Windows Subsystem for Linux (WSL) 2 introduces a significant architectural change as it is a full Linux kernel built by Microsoft, allowing Linux containers to run natively without emulation.
- <https://www.microsoft.com/en-us/software-download/windows10>

The screenshot shows a web browser displaying the Docker Docs website at <https://docs.docker.com/docker-for-windows/wsl/>. The page title is "Docker Desktop WSL 2 backend". The left sidebar contains a navigation menu with links such as "Install Docker Desktop for Windows", "User manual", "Deploy on Kubernetes", "Networking", "Logs and troubleshooting", "FAQs", "Stable release notes", "Edge release notes", "Docker Desktop WSL 2 backend" (which is highlighted), "Dashboard", and "Open source licensing". The main content area starts with a section titled "Prerequisites" which lists three steps: 1. Install Windows 10, version 1903 or higher. 2. Enable WSL 2 feature on Windows. 3. Download and install the Linux kernel update package. Below this is a section titled "Best practices" with a single bullet point about bind-mounting files. On the right side of the page, there are options to "Edit this page", "Request docs changes", and a "On this page" sidebar listing links to "Prerequisites", "Best practices", "Download", "Install", "Develop with Docker and WSL 2", and "Feedback".

<https://docs.docker.com/docker-for-windows/wsl/>

Try It!

Basic Docker Commands

- *docker ps*
- *docker ps -a*
- *docker images*
- *docker run --name some-nginx -p 8080:80 -d nginx*
- *docker ps*

docker ps					
CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
NAMES					
79ed823b6a79	nginx	"/docker-entrypoint..."	10 minutes ago	Up 10 minutes	0.0.0.0:8080->80/tcp
some-nginx					

- Open browser and try <http://localhost:8080>
- *docker images*

Basic Docker Commands

Commands:

```
attach          Attach to a running container
build           Build an image from a Dockerfile
commit          Create a new image from a container's changes
cp              Copy files/folders between a container and the local filesystem
create          Create a new container
diff            Inspect changes on a container's filesystem
events          Get real time events from the server
exec            Run a command in a running container
export          Export a container's filesystem as a tar archive
history         Show the history of an image
images          List images
import          Import the contents from a tarball to create a filesystem image
info             Display system-wide information
inspect         Return low-level information on a container, image or task
kill             Kill one or more running container
load             Load an image from a tar archive or STDIN
login            Log in to a Docker registry.
logout           Log out from a Docker registry.
logs             Fetch the logs of a container
network          Manage Docker networks
node             Manage Docker Swarm nodes
pause            Pause all processes within one or more containers
port             List port mappings or a specific mapping for the container
ps               List containers
pull             Pull an image or a repository from a registry
push             Push an image or a repository to a registry
rename           Rename a container
restart          Restart a container
rm               Remove one or more containers
rmi              Remove one or more images
run              Run a command in a new container
save             Save one or more images to a tar archive <streamed to STDOUT by default>
search            Search the Docker Hub for images
service           Manage Docker services
start            Start one or more stopped containers
stats             Display a live stream of container(s) resource usage statistics
stop              Stop one or more running containers
swarm             Manage Docker Swarm
tag               Tag an image into a repository
top                Display the running processes of a container
unpause           Unpause all processes within one or more containers
update            Update configuration of one or more containers
version           Show the Docker version information
volume            Manage Docker volumes
wait              Block until a container stops, then print its exit code
```

- docker ps
- docker ps -a
- docker run
- docker start
- docker stop
- docker logs
- docker pull
- docker images
- docker rm
- docker rmi
- docker info
- docker inspect

Basic Docker Commands (con't)

- *docker images*
- *docker logs <containerId>*
- *docker logs -f <containerId>*
- *docker stop <containerId>*
- *docker ps -a*
- *docker start <containerId>*
- *docker inspect <containerId>*
- **Limit a container's resources** <https://docs.docker.com/engine/admin/resource_constraints/>

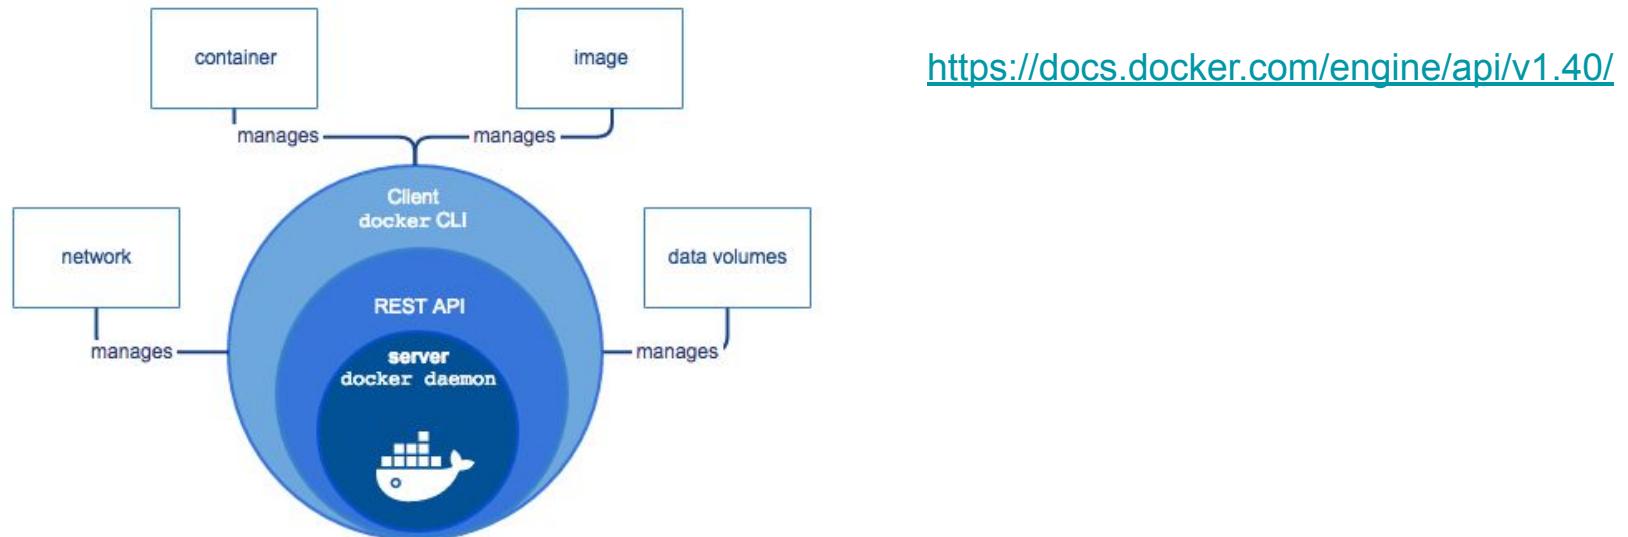
Note:

- Can use **container name** instead of **containerId**
- *docker <command> --help*
- *Docker <command> <subcommand> --help*

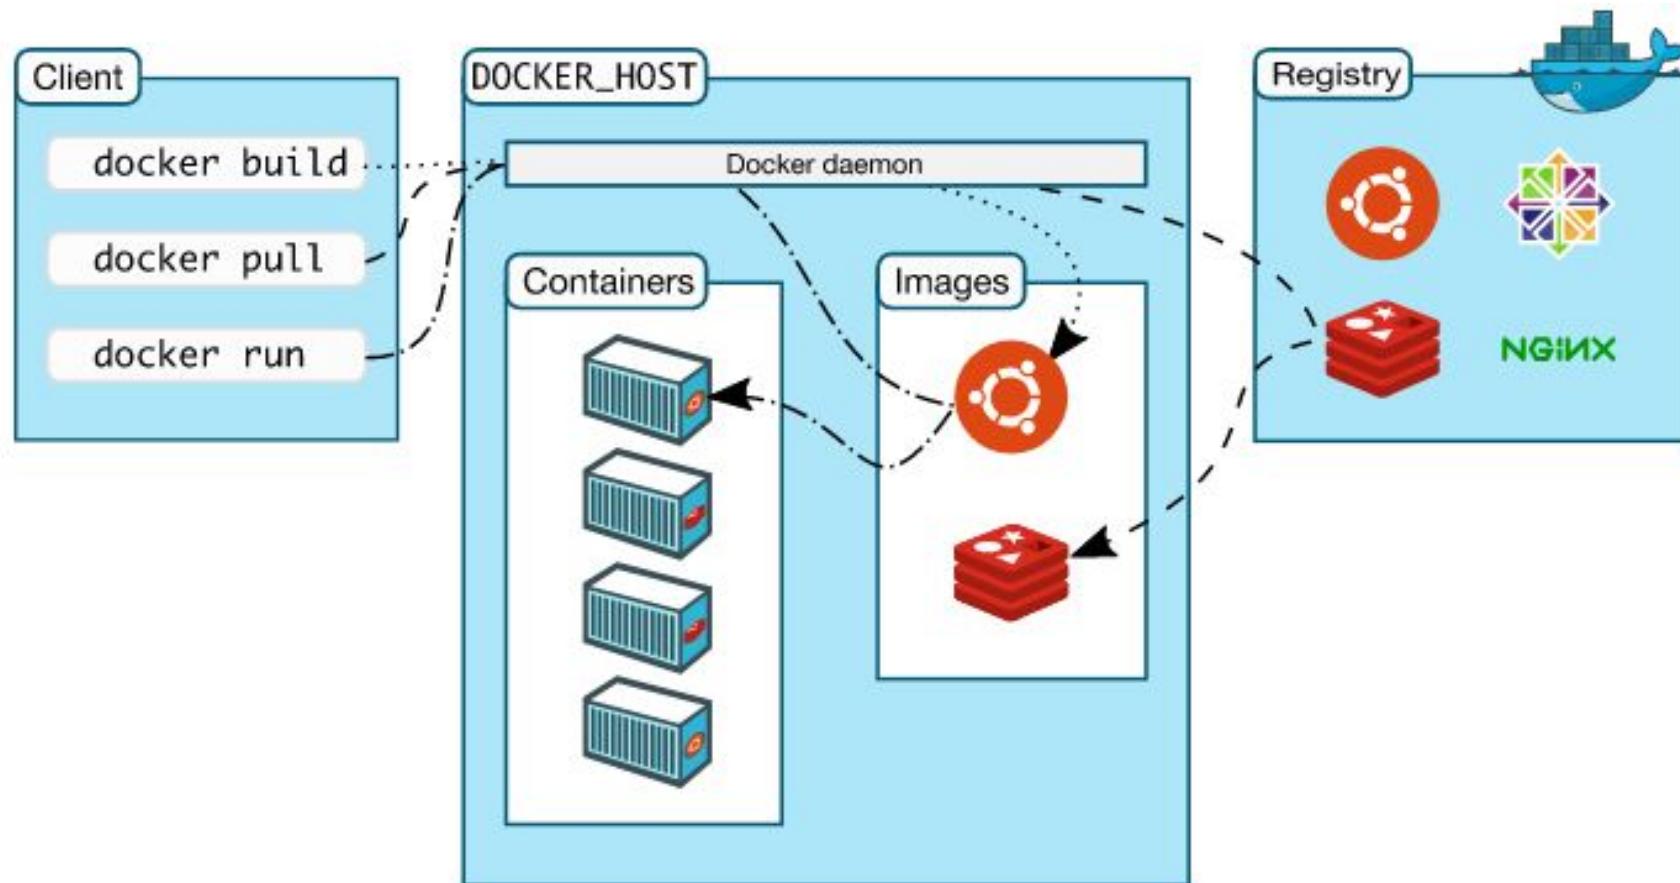
Docker Engine

Docker Engine is a **client-server application** with these major components:

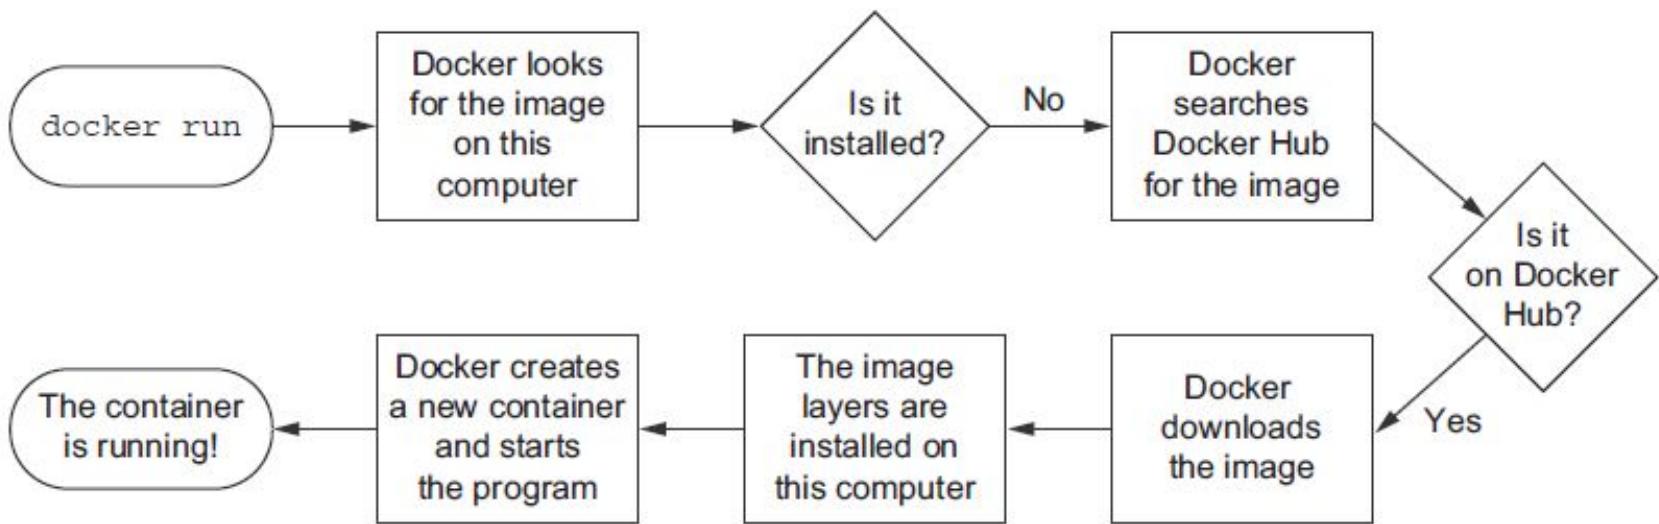
- A server which is a type of long-running program called **a daemon process** (the dockerd command).
- A REST API which specifies interfaces that programs can use to talk to the daemon and instruct it what to do.
- A command line interface (CLI) client (the docker command).
- `curl --unix-socket /var/run/docker.sock http://localhost/v1.40/containers/json`



Docker Architecture



Flow of docker run command



Inside Docker Desktop

Mac

- *screen ~/Library/Containers/com.docker.docker/Data/vms/0/tty*
- *cd /var/lib/docker*

Windows

- *docker run -it --privileged --pid=host debian nsenter -t 1 -m -u -i sh*
- *cd /var/lib/docker*

Container List

✓ docker ps	CONTAINER ID	IMAGE	COMMAND NAMES	CREATED	STATUS	PORTS
	6fb1c7b7bd84	portainer/portainer	"portainer"	About an hour ago	Up About an hour	0.0.0.0:9001->9000/tcp
	9adb7c194b0e	docker.elastic.co/kibana/kibana:6.3.2	"portainer"	11 days ago	Up 11 days	0.0.0.0:5601->5601/tcp
	639718c9aa3e	bitnami/redis-sentinel:latest	"kibana"	3 months ago	Up 3 weeks	0.0.0.0:26380->26379/tcp
	fd00c9463a41	bitnami/redis-sentinel:latest	"bitnami_redis-sentinel_2"	3 months ago	Up 3 weeks	0.0.0.0:26381->26379/tcp
	76bcc50ac645	bitnami/redis-sentinel:latest	"bitnami_redis-sentinel_3"	3 months ago	Up 3 weeks	0.0.0.0:26379->26379/tcp
	b5c69f162304	bitnami/redis:latest	"bitnami_redis-sentinel_1"	3 months ago	Up 3 weeks	0.0.0.0:32769->6379/tcp
	99303b866e8c	bitnami/redis:latest	"bitnami_redis-slave_1"	3 months ago	Up 3 weeks	0.0.0.0:32768->6379/tcp
	e7cd519a70a9	prom/prometheus:v2.20.0	"bitnami_redis_1"	3 months ago	Up 3 weeks	9090/tcp
	7361c250a9ea	prom/alertmanager:v0.21.0	"prometheus"	3 months ago	Up 3 weeks	9093/tcp
	304fe3414086	gcr.io/cadvisor/cadvisor:v0.37.0	"alertmanager"	3 months ago	Up 3 weeks (healthy)	8080/tcp
	2adcef6ed50e .0.0:9090->9091->9090-9091/tcp, 0.0.0.0:9093->9093/tcp	stefanprodan/caddy	"cadvisor"	3 months ago	Up 3 weeks	0.0.0.0:3000->3000/tcp, 0.0
	2f83908fa48e	grafana/grafana:7.1.1	"caddy"	3 months ago	Up 3 weeks	3000/tcp
	5a63958e75df	prom/node-exporter:v1.0.1	"grafana"	3 months ago	Up 3 weeks	9100/tcp
	65e9e2ec7ffe	prom/pushgateway:v1.2.0	"nodeexporter"	3 months ago	Up 3 weeks	9091/tcp
	c21d70f731d5	elkozmon/zoonavigator:latest	"pushgateway"	3 months ago	Up 3 weeks	0.0.0.0:9000->9000/tcp
	74c1f48fdc44	anutech2001/reader	"zoonavigator"	4 months ago	Up 3 weeks (healthy)	0.0.0.0:8010->8010/tcp
	d4d277e37279 8/tcp	anutech2001/bookstore	"training_reader_1"	5 months ago	Up 3 weeks	0.0.0.0:8011->8011/tcp, 888
	48743fb7b571 1/tcp	openzipkin/zipkin	"training_bookstore_1"	5 months ago	Up 3 weeks	9410/tcp, 0.0.0.0:9411->941
			"busybox/sh run.sh"	5 months ago	Up 3 weeks	
			"training_zipkin_1"			

Container Process List

- \$ top

PID	PPID	USER	STAT	VSZ	%VSZ	CPU	%CPU	COMMAND
9177	7901	root	S	205m	3%	7	3%	/usr/bin/cadvisor -logtostderr
1640	1632	root	S	2297m	28%	4	0%	/usr/local/bin/dockerd -H unix:///
8958	7714	65534	S	702m	9%	6	0%	/bin/node_exporter --path.procfs=/
7514	7369	1000	S	6412m	78%	1	0%	/usr/local/openjdk-11/bin/java -Du
7868	7627	65534	S	1438m	18%	5	0%	/bin/prometheus --config.file=/etc
1654	1640	root	S	2033m	25%	3	0%	containerd --config /var/run/docke
1090	1	root	S	142m	2%	1	0%	/usr/bin/containerd
484	1	root	S	104m	1%	5	0%	/usr/bin/memlogd -fd-log 3 -fd-que
8855	7686	1001	S	52792	1%	0	0%	redis-sentinel 0.0.0.0:26379 [sent
7199	6986	1001	S	52792	1%	1	0%	redis-sentinel 0.0.0.0:26379 [sent
7217	7084	1001	S	52792	1%	6	0%	redis-sentinel 0.0.0.0:26379 [sent
7369	1654	root	S	106m	1%	5	0%	containerd-shim -namespace moby -w
7169	6844	472	S	740m	9%	6	0%	grafana-server --homepath=/usr/sha
7302	7124	1001	S	138m	2%	3	0%	redis-server 0.0.0.0:6379
7901	1654	root	S	106m	1%	7	0%	containerd-shim -namespace moby -w
7473	7276	root	S	5594m	68%	0	0%	{java} /usr/lib/jvm/java-8-openjdk
7182	6994	root	S	5458m	67%	2	0%	

Docker Registry

Docker Hub (Public Registry)

[Docker Hub](https://hub.docker.com/) (<https://hub.docker.com/>) provides a free-to-use, hosted Registry, plus additional features (organization accounts, automated builds, and more).

The screenshot shows the Docker Hub homepage with a blue header bar. On the left, there's a search bar with placeholder text "Search for great content (e.g., mysql)". On the right, there are links for "Explore", "Pricing", "Sign In", and a prominent "Sign Up" button.

Build and Ship any Application Anywhere

Docker Hub is the world's easiest way to create, manage, and deliver your teams' container applications.

Sign Up Today

Already have an account? [Sign In](#)

Docker ID

Email

Password

Send me occasional product updates and announcements.

I'm not a robot reCAPTCHA

[Sign Up](#)

By creating an account, you agree to the [Terms of Service](#), [Privacy Policy](#), and [Data Processing Terms](#).

Try It!

Private Docker Registry

For faster access and shared an images, create our private registry.

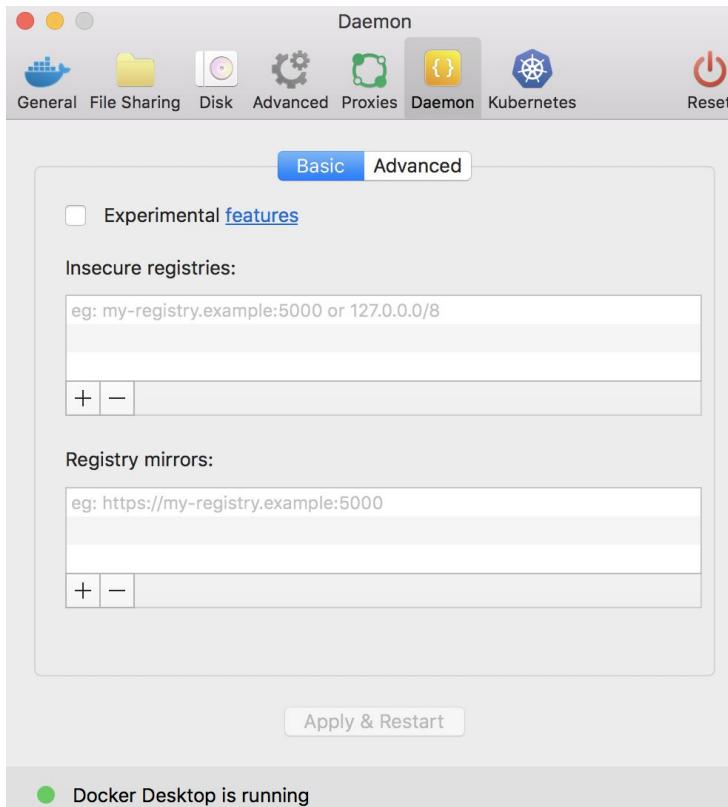
- Run your private registry
docker run -d -p 5000:5000 --restart=always --name registry registry:2
- Pull some image from the Docker Hub
docker pull hello-world
- Tag the image so that it points to your registry
docker tag hello-world localhost:5000/myhelloworld
- Push it
docker push localhost:5000/myhelloworld
- Pull it back (from another docker machine)
docker pull localhost:5000/myhelloworld
- List Catalog
http://<registryHost>:5000/v2/_catalog
- See Tags of an Image, for example from “myhelloworld” image
http://localhost:5000/v2/myhelloworld/tags/list

Try It!

Upload Image to Docker Hub

- Step to Upload Image to Docker Hub
 - Register Free Docker Hub accounts (<https://hub.docker.com/>)
 - *docker tag <image name>:<version> <docker hub user>/<image name>:<version>*
 - *docker login –u <docker hub user> –p <docker hub password>*
 - *docker push <docker hub user>/<image>:<version>*
- Example
 - *docker tag hello-world <docker hub user>/myhello-world*
 - *docker login –u <docker hub user>*
 - *docker push <docker hub user>/myhello-world*
 - Delete image
docker rmi <image name>

Setting External Private Registry



Docker Engine
v19.03.13

Configure the Docker daemon by typing a json Docker daemon [configuration](#) file.
This can prevent Docker from starting. Use at your own risk!

```
{  
  "registry-mirrors": [],  
  "insecure-registries": [],  
  "debug": true,  
  "experimental": false  
}
```

Advantage of Private Registry

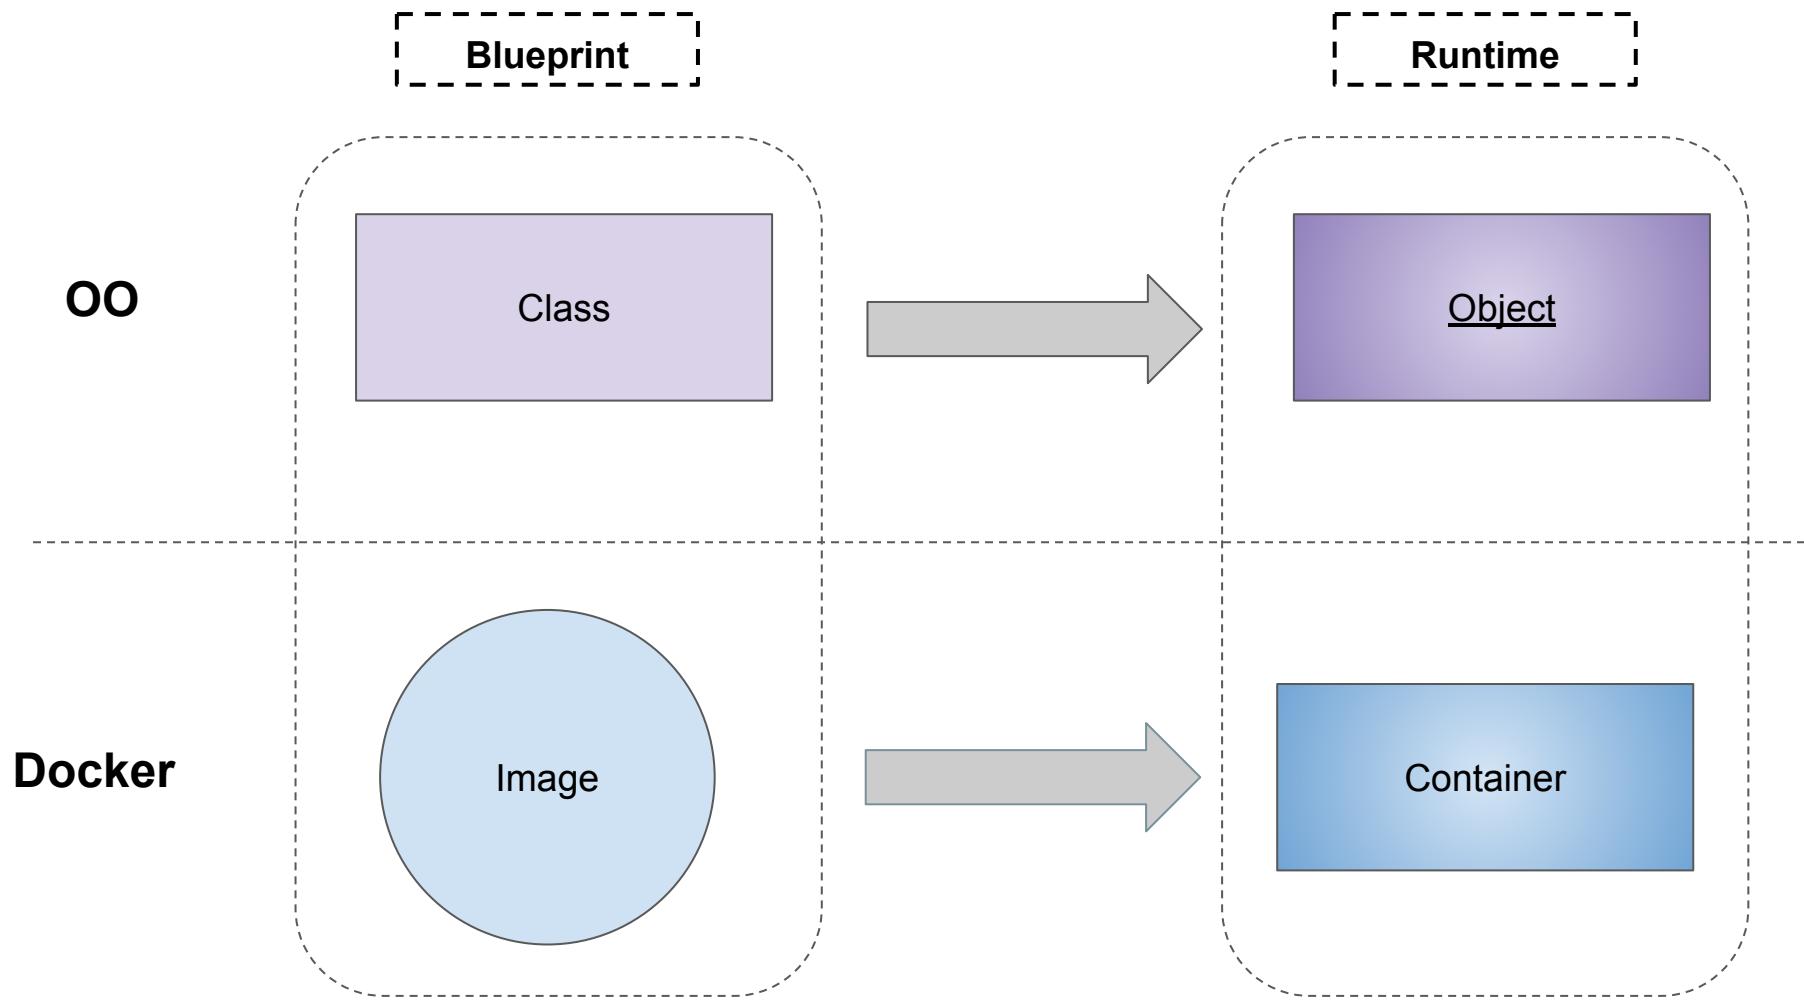
- Speed
- Security
- Reliability
- Regulation
- Cost



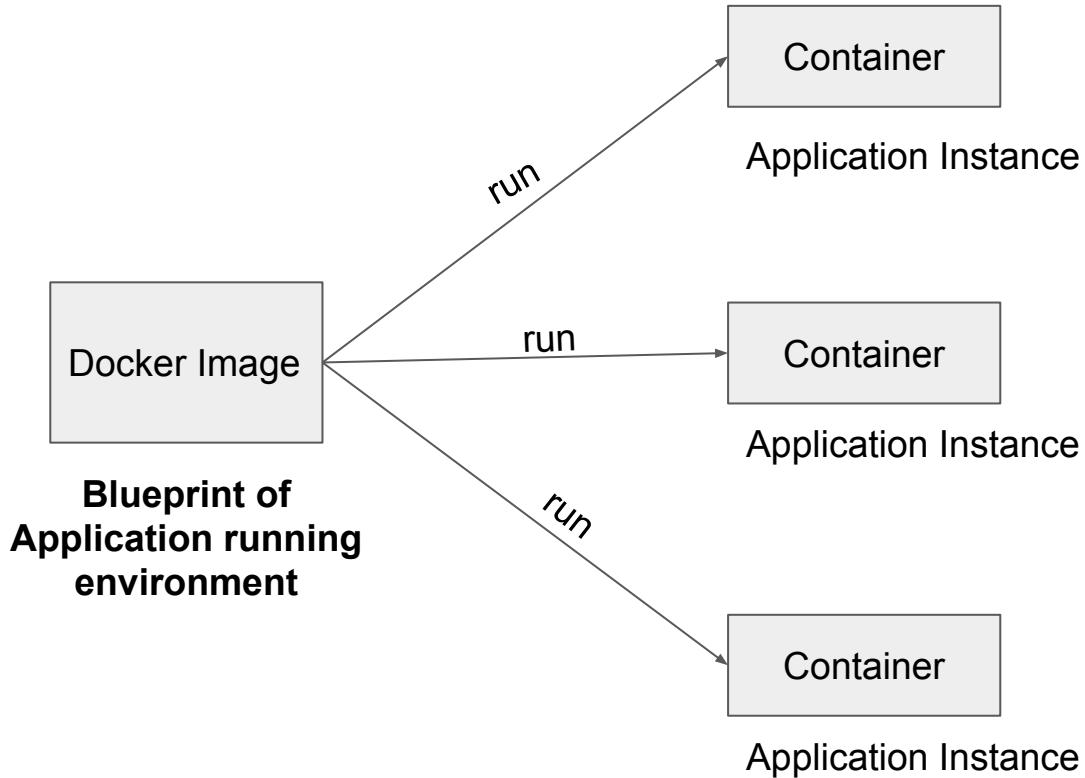
opensuse/portus

Docker Images

Compare with Object-Oriented



Docker Image



Docker Images and Layers

91e54dfb1179	0 B
d74508fb6632	1.895 KB
c22013c84729	194.5 KB
d3a1f33e8a5a	188.1 MB
ubuntu:15.04	
Image	

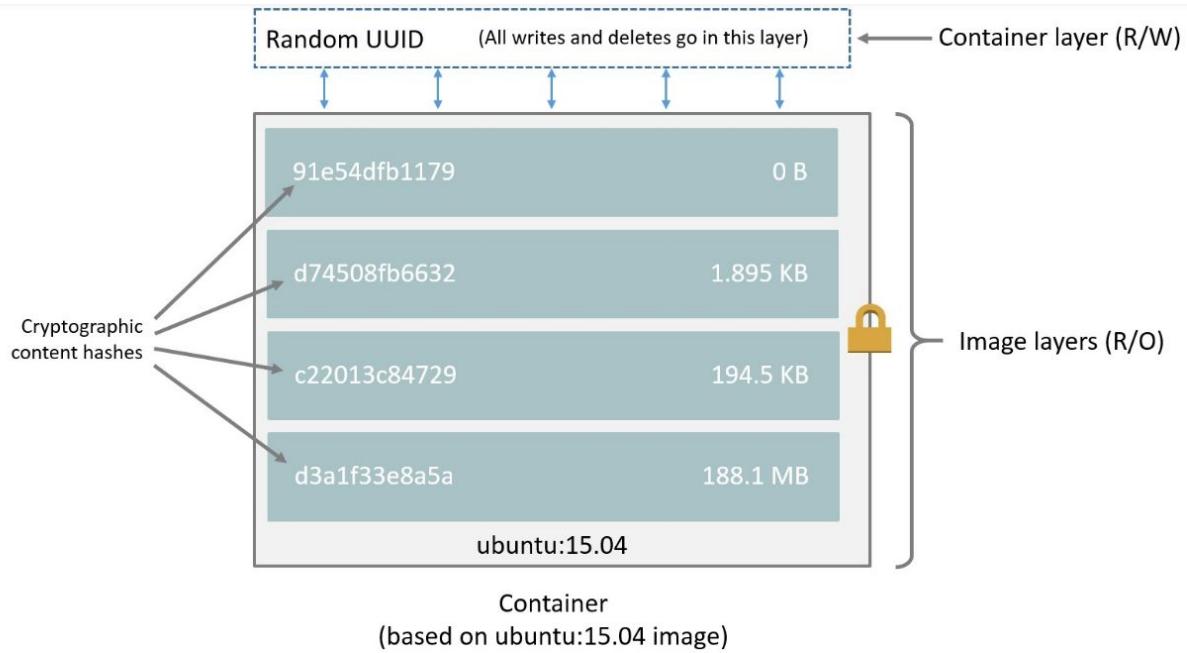
- Each Docker image references a **list of read-only layers** that represent filesystem differences.
- The **Docker storage driver** is responsible for stacking these layers and providing a single unified view.

Container's Key Technologies

Two key technologies behind Docker image and container management.

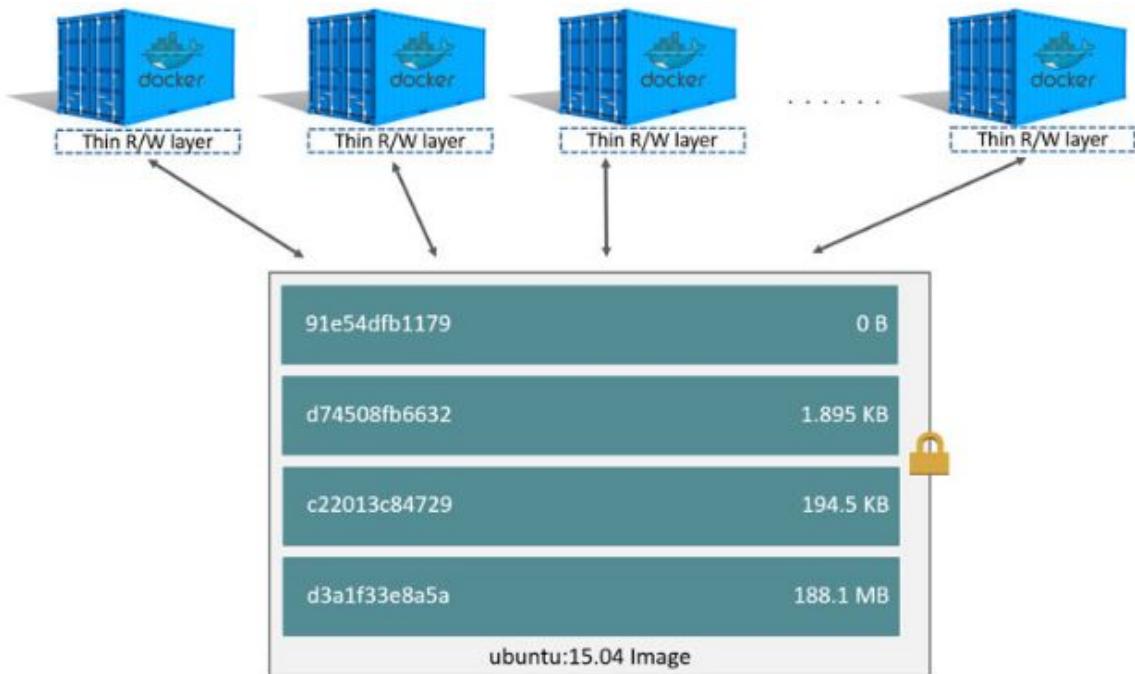
- ***Stackable image layers***
- ***The copy-on-write (CoW) strategy:***
 - System processes that need the same data share the same instance of that data.
 - If one process needs to modify or write to the data, only then does the operating system make a copy of the data for that process to use and only the process that needs to write has access to the data copy.
- ***These strategies optimizes both image disk space usage and the performance of container start times***

Container and Layers



- New R/W layer was created when the container was created.
- All changes are written to this layer.
- When the container is deleted the writable layer is also deleted.
- Since Docker 1.10

Container and Layers (cont')



- Multiple containers can share access to the same underlying image and yet have their own data state.
- Sharing is a good way to optimize resources.

Dockerfile of jdk:8

```
1 FROM buildpack-deps:jessie-scm
2
3 RUN apt-get update && apt-get install -y --no-install-recommends \
4     bzip2 \
5     unzip \
6     xz-utils \
7     && rm -rf /var/lib/apt/lists/*
8 RUN echo 'deb http://deb.debian.org/debian jessie-backports main' > /etc/apt/sources.list.d/jessie-backports.list
9 ENV LANG C.UTF-8
10 RUN { \
11     echo '#!/bin/sh'; \
12     echo 'set -e'; \
13     echo; \
14     echo 'dirname "$(dirname "$(readlink -f "$(which javac || which java)")")"'; \
15 } > /usr/local/bin/docker-java-home \
16 && chmod +x /usr/local/bin/docker-java-home
17
18 ENV JAVA_HOME /usr/lib/jvm/java-8-openjdk-amd64
19
20 ENV JAVA_VERSION 8u111
21 ENV JAVA_DEBIAN_VERSION 8u111-b14-2~bpo8+1
22 ENV CA_CERTIFICATES_JAVA_VERSION 20140324
23
24 RUN set -x \
25     && apt-get update \
26     && apt-get install -y \
27         openjdk-8-jdk="$JAVA_DEBIAN_VERSION" \
28         ca-certificates-java="$CA_CERTIFICATES_JAVA_VERSION" \
29     && rm -rf /var/lib/apt/lists/* \
30     && [ "$JAVA_HOME" = "$(docker-java-home)" ]
31 RUN /var/lib/dpkg/info/ca-certificates-java.postinst configure
```

Docker Image of java:8

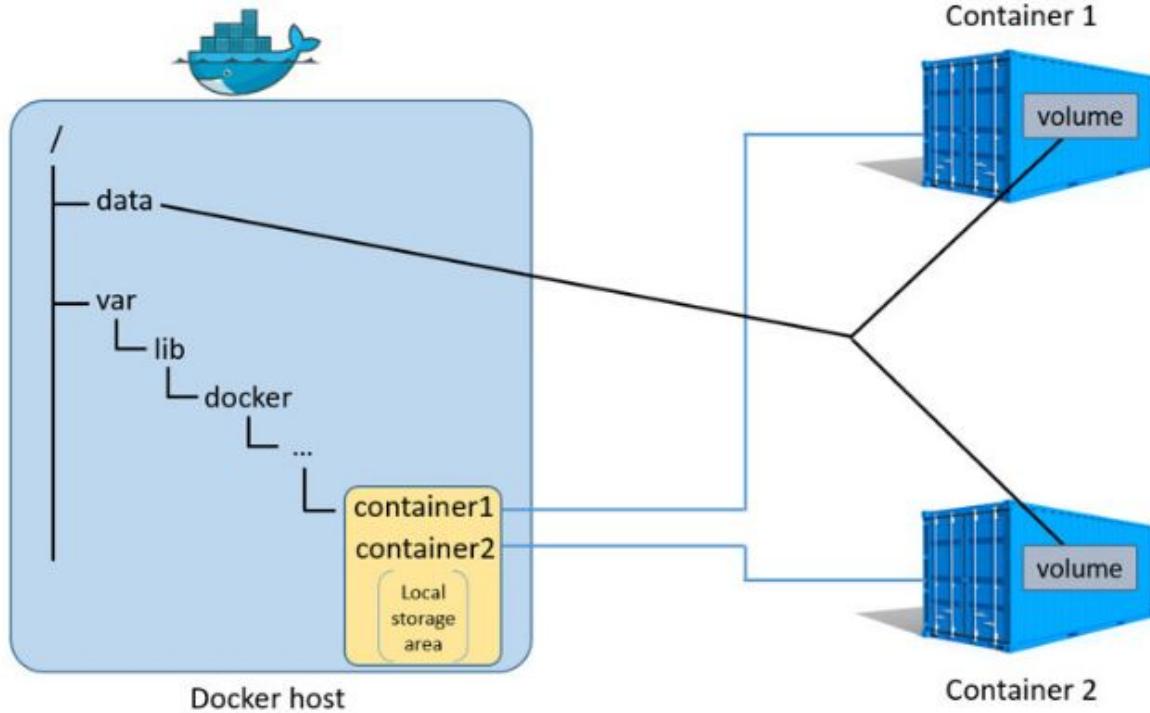
Image

192.168.99.117:5000/java:8

Image	Cmd	Size
54970924248	/var/lib/dpkg/info/ca-certificates-java.postinst configure	285 KB
71b6ed1bb5e	/bin/sh -c set -x && apt-get update && apt-get install -y openjdk-8-jdk="\$JAVA_DEBIAN_VERSION" ca-certificates-java="\$CA_CERTIFICATES_JAVA_VERSION" && rm -rf /var/lib/apt/lists/* && ["\$JAVA_HOME" = "\$(docker-java-home)"]	131.6 MB
c2e26dad9eb	/bin/sh -c #(nop) ENV CA_CERTIFICATES_JAVA_VERSION=20140324	0 B
3305d7e2c8a	/bin/sh -c #(nop) ENV JAVA_DEBIAN_VERSION=8u102-b14.1-1~bpo8+1	0 B
a001ffeb3ab	/bin/sh -c #(nop) ENV JAVA_VERSION=8u102	0 B
a76495019d4	/bin/sh -c #(nop) ENV JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64	0 B
031d51a8c22	/bin/sh -c { echo '#!/bin/sh'; echo 'set -e'; echo; echo 'dirname "\$(dirname "\$(readlink -f "\$(which javac which java)")")"'; } > /usr/local/bin/docker-java-home && chmod +x /usr/local/bin/docker-java-home	240 B
a43d5ea7d36	/bin/sh -c #(nop) ENV LANG=C.UTF-8	0 B
6141a4dc4ed	/bin/sh -c echo 'deb http://httpredir.debian.org/debian jessie-backports main' > /etc/apt/sources.list.d/jessie-backports.list	218 B
eecb52a8466	/bin/sh -c apt-get update && apt-get install -y --no-install-recommends bzip2 unzip xz-utils && rm -rf /var/lib/apt/lists/*	599 KB
5d459c46dd1	/bin/sh -c apt-get update && apt-get install -y --no-install-recommends bzr git mercurial openssh-client subversion procps && rm -rf /var/lib/apt/lists/*	43.2 MB
488a5a26b89	/bin/sh -c apt-get update && apt-get install -y --no-install-recommends ca-certificates curl wget && rm -rf /var/lib/apt/lists/*	19.2 MB
51dd77e0947	/bin/sh -c #(nop) CMD ["/bin/bash"]	0 B
57f55a8a684	/bin/sh -c #(nop) ADD file:f2453b914e7e026efd39c6321c7b14509b6d09dd3cf5567a8f6bd38466e06954 in /	52.5 MB

Total size: 247.3 MB

The storage driver and Data volumes



- A single Docker host running two containers.
- Each container own address space within the Docker host's local storage area (`/var/lib/docker/...`)
- A single shared data volume located at `/data` on the Docker host

Portainer

An open-source toolset that allows you to easily build and manage containers in Docker, Swarm, Kubernetes and Azure ACI.

```
$ docker volume create portainer_data
```

Win: docker run -d -p 8000:8000 -p 9000:9000 --name=portainer --restart=always -v /var/run/docker.sock:/var/run/docker.sock -v portainer_data:/data portainer/portainer-ce

Mac: docker run -d -p 9001:9000 --name portainer --restart always -v /var/run/docker.sock:/var/run/docker.sock -v portainer_data:/data portainer/portainer

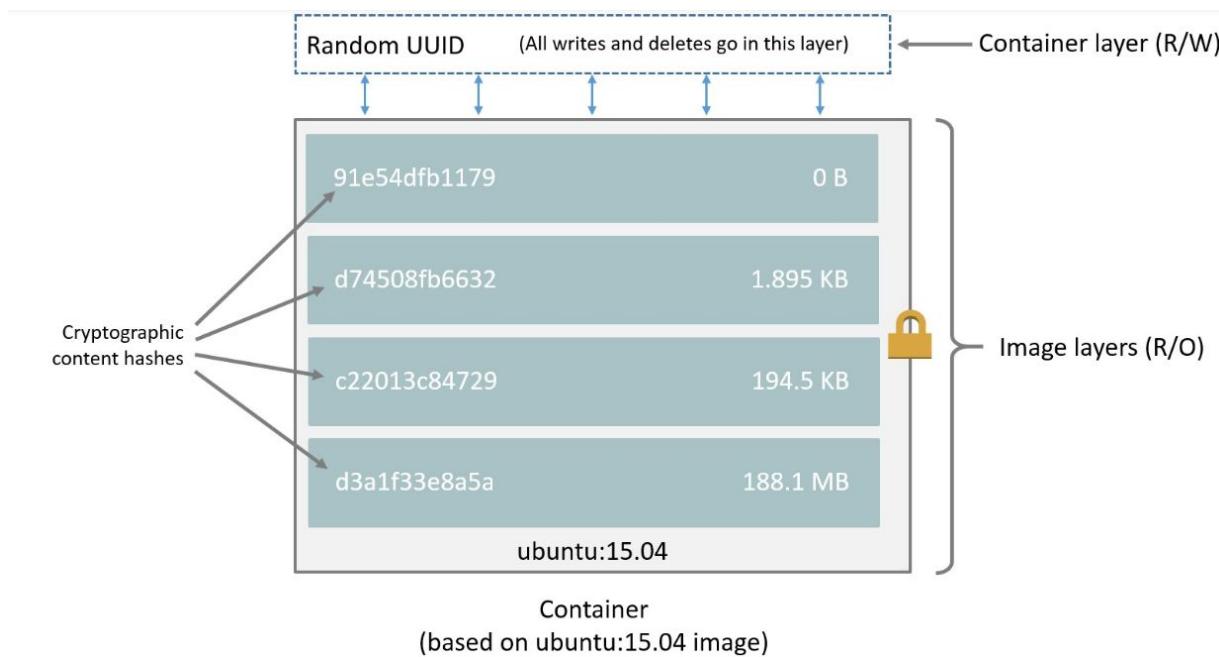
The screenshot shows the Portainer.io dashboard interface. On the left is a dark sidebar with a navigation menu. The main area is titled "Dashboard" and displays "Endpoint summary". It includes a "Endpoint info" section with details like "local", "8", "9.9 GB - Standalone 19.03.13", and a URL of "/var/run/docker.sock". Below this are four large cards: "Stacks" (0), "Containers" (16, with 0 healthy, 4 running, 0 unhealthy, and 12 stopped), "Images" (29, 3.4 GB), and "Volumes" (3). At the bottom is a card for "Networks" (3).

Docker Volume

Manage data in Docker

By default all files created inside a container are **stored on a writable container layer**.

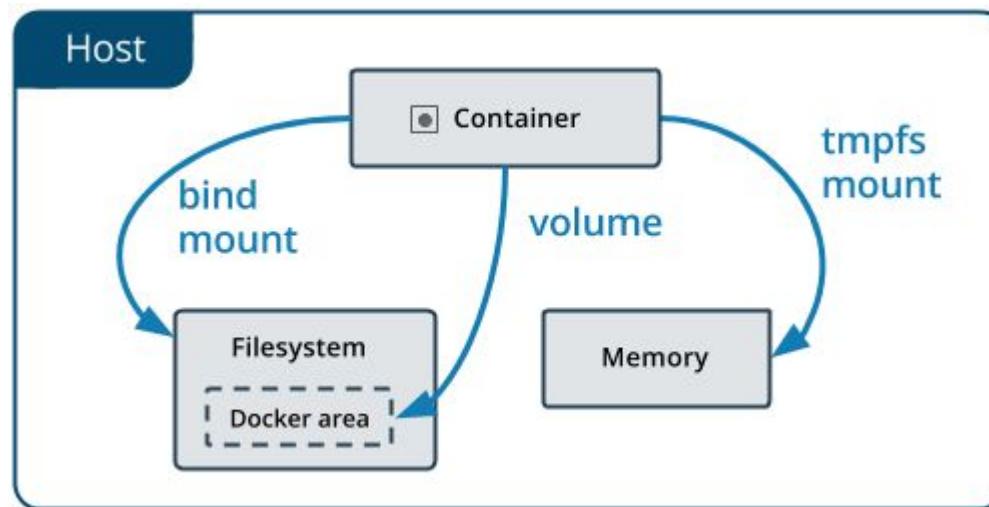
- The data doesn't persist when that container no longer exists.
- This extra abstraction reduces performance.



Different Types of Manage Data

Two options for containers to store files in the host machine.

- **Volumes** are stored in a part of the host filesystem which is managed by Docker (/var/lib/docker/volumes/)
- **Bind mounts** may be stored anywhere on the host system.



Note: No matter which type of mount you choose to use, **the data looks the same from within the container.**

Docker run with Bind mounts

- *docker run --name some-nginx -v /c/Users/<user>/nginx:/usr/share/nginx/html:ro -p 8080:80 -d nginx*
- Locate a volume
 - *docker inspect*

```
"Mounts": [
  {
    "Type": "volume",
    "Name": "1fb394327342b55036eea22e3b24c2e4a45a352ba673ab0696af01a103c274ba",
    "Source": "/var/lib/docker/volumes/1fb394327342b55036eea22e3b24c2e4a45a352ba673ab0696af01a103c274ba/_data",
    "Destination": "/tmp/a",
    "Driver": "local",
    "Mode": "",
    "RW": true,
    "Propagation": ""
  }
],
```

- <https://docs.docker.com/storage/volumes/>

Docker run with Volume (anonymous volume and named volume)

Create a volume explicitly

```
docker volume create awesome
```

```
docker run --rm -v /foo -v awesome:/bar busybox top
```

Or just:

```
docker run --rm -v /foo -v awesome:/bar busybox top
```

Inspect Volume

```
docker volume inspect <volume name>
```

Remove volumes

```
docker volume prune
```

```
docker volume rm <volume name>
```

Good use cases for volumes and Bind mounts

Volumes

- Sharing data among multiple running containers.
- When you want to store your container's data on a remote host or a cloud provider, rather than locally.
- When you need to back up, restore, or migrate data from one Docker host to another, volumes are a better choice.
- When your application requires high-performance I/O on Docker Desktop.

Bind mounts

- Sharing configuration files from the host machine to containers.
- Sharing source code or build artifacts between a development environment on the Docker host and a container.

Advanced Topics

[Use a volume driver](#)

[Backup, restore, or migrate data volumes](#)

[Docker storage drivers](#)

Docker Network

Network driver summary

- **User-defined bridge networks** are best when you need multiple containers to communicate on the same Docker host.
- **Host networks** are best when the network stack should not be isolated from the Docker host, but you want other aspects of the container to be isolated.
- **Overlay networks** are best when you need containers running on different Docker hosts to communicate, or when multiple applications work together using swarm services.
- **Macvlan networks** are best when you are migrating from a VM setup or need your containers to look like physical hosts on your network, each with a unique MAC address.
- **Third-party network plugins** allow you to integrate Docker with specialized network stacks.

Bridge Networks

- The default network driver. If you don't specify a driver.
- A software bridge which allows containers connected to the same bridge network to communicate, while providing isolation from containers which are not connected to that bridge network.
- Bridge networks apply to containers running on the same Docker daemon host
- When you start Docker, a default bridge network (also called “bridge”) is created automatically
- Basic network commands
 - docker network ls
 - docker network inspect <network name>
 - docker network create --driver <driver type> <network name>
 - docker network rm <network name>

Noted: driver type can be “bridge”, “host”, “overlay”, “macvlan”, or “none”.

Bridge Network Examples

```
$ docker run -dit --name alpine1 alpine ash
```

```
$ docker run -dit --name alpine2 alpine ash
```

```
$ docker network inspect bridge #note ip of alpine2
```

```
$ docker attach alpine1
```

```
$ ping -c 2 <alpine2 ip>
```

```
$ ping -c 2 alpine2
```

What happened in the last command?

User-defined Bridge Networks

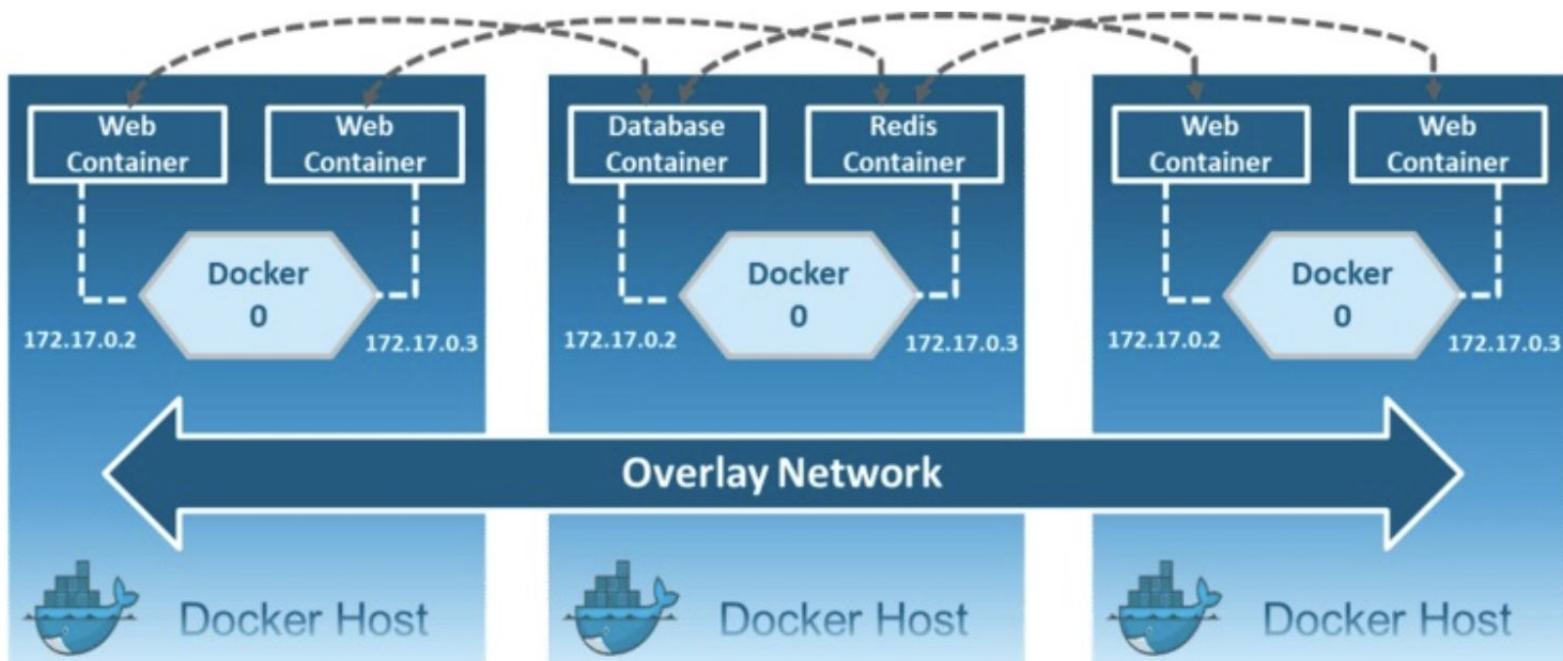
- Try these commands
 - `$ docker network create my-net`
 - `$ docker run -dit --network my-net --name alpine1 alpine ash`
 - `$ docker run -dit --network my-net --name alpine2 alpine ash`
 - `$ docker network inspect my-net --- noted an ip of alpine2`
 - `$ docker attach alpine1`
 - `$ ping -c 2 <alpine2 ip>`
 - `$ ping -c 2 alpine2`
- What happened in the last command?
- `$ docker network rm my-net`
- `$ docker rm $(docker ps -a -q -f status=exited)`

User-defined Bridge Networks vs Default Bridge

- Differences Between User-defined Bridges and The Default Bridge
 - User-defined bridges provide automatic DNS resolution between containers.
 - User-defined bridges provide better isolation.
 - Containers can be attached and detached from user-defined networks on the fly.
 - Each user-defined network creates a configurable bridge.
 - Linked containers on the default bridge network share environment variables.
- <https://docs.docker.com/network/bridge/#differences-between-user-defined-bridges-and-the-default-bridge>
- https://hub.docker.com/_/busybox
- the special DNS name `host.docker.internal`, which resolves to the internal IP address used by the host.

Overlay Networks

- The overlay network driver creates a distributed network among multiple Docker daemon hosts.
- Creates an internal private network that spans across all the nodes participating in the swarm cluster.
- Overlay networks facilitate communication between a swarm service and a standalone container, or between two standalone containers on different Docker Daemons.



Create Overlay Network

- Create an overlay network

```
$ docker network create -d overlay my-overlay
```

- To create an overlay network which can be used by swarm services or standalone containers to communicate with other standalone containers running on other Docker daemons, add the `--attachable` flag:

```
$ docker network create -d overlay --attachable my-attachable-overlay
```

- You can specify the IP address range, subnet, gateway, and other options. See `docker network create --help` for details.
- To encrypt application data as well, add `--opt encrypted` when creating the overlay network. This enables IPSEC encryption at the level of the vxlan.**Overlay network encryption is not supported on Windows nodes.**

```
$ docker network create --opt encrypted --driver overlay --attachable  
my-attachable-multi-host-network
```

Docker Cheat Sheet

General Usage

Start a container in background
`$ docker run -d jenkins`

Start an interactive container
`$ docker run -it ubuntu bash`

Start a container automatically removed on stop
`$ docker run -rm ubuntu bash`

Export port from a container
`$ docker run -p 80:80 -d nginx`

Start a named container
`$ docker run --name mydb redis`

Restart a stopped container
`$ docker start mydb`

Stop a container
`$ docker stop mydb`

Add metadata to container
`$ docker run -d \label=traefik.backend=jenkins jenkins`

Build Images

Build an image from Dockerfile in current directory
`$ docker build --tag myimage .`

Force rebuild of Docker image
`$ docker build --no-cache .`

Convert a container to image
`$ docker commit c7337 myimage`

Remove all unused images
`$ docker rmi $(docker images \ -q -f "dangling=true")`

Debug

Run another process in running container
`$ docker exec -it c7337 bash`

Show live logs of running daemon container
`$ docker logs -f c7337`

Show exposed ports of a container
`$ docker port c7337`

Manage Containers

List running containers
`$ docker ps`

List all containers (running & stopped)
`$ docker ps -a`

Inspect containers metadatas
`$ docker inspect c7337`

List local available images
`$ docker images`

Delete all stopped containers
`$ docker rm $(docker ps --filter status=exited -q)`

List all containers with a specific label
`$ docker ps --filter label=traefik.backend`

Query a specific metadata of a running container
`$ docker inspect -f '{{ .NetworkSettings.IPAddress }}' c7337`



Volumes

Create a local volume
`$ docker volume create --name myvol`

Mounting a volume on container start
`$ docker run -v myvol:/data redis`

Destroy a volume
`$ docker volume rm myvol`

List volumes
`$ docker volume ls`

Create a local network
`$ docker network create mynet`

Attach a container to a network on start
`$ docker run -d --net mynet redis`

Connect a running container from a network
`$ docker network connect mynet c7337`

Disconnect container to a network
`$ docker network disconnect mynet c7337`

Legend

Image name
`redis, jenkins, nginx`

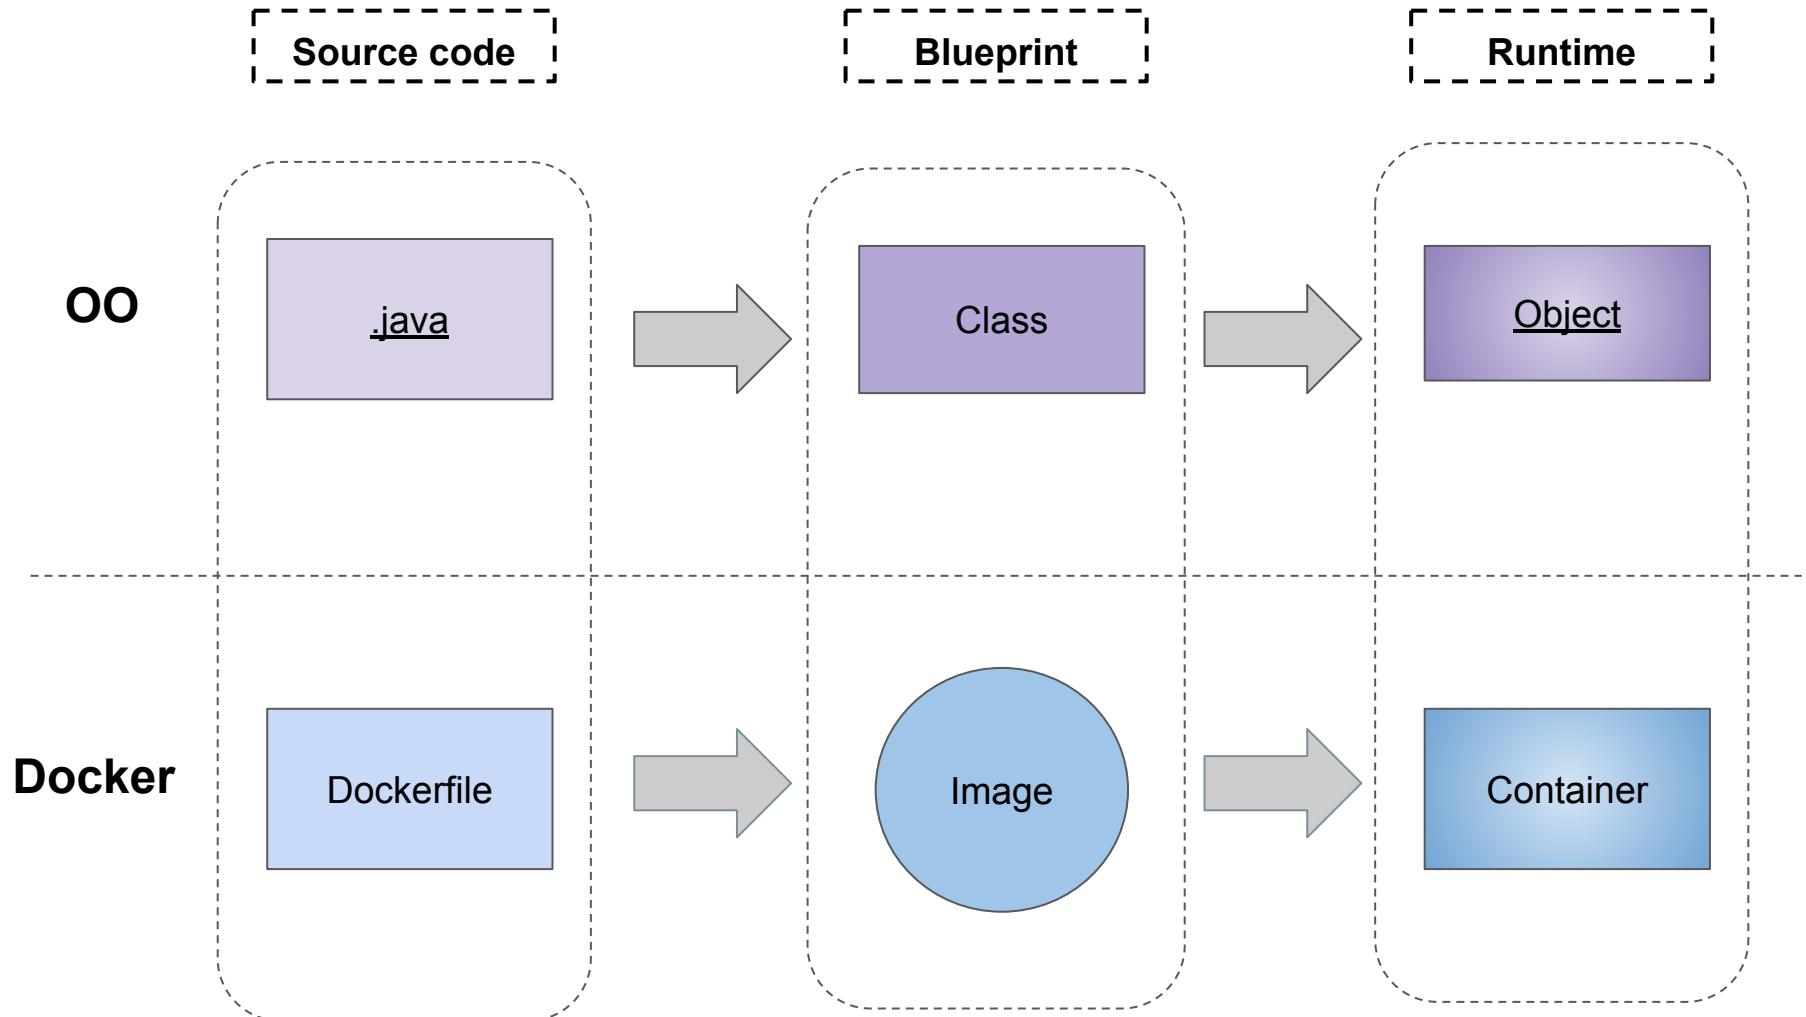
Container name or commit ID
`mydb #name
c7337 #commit id`

Docker Build (Dockerfile)

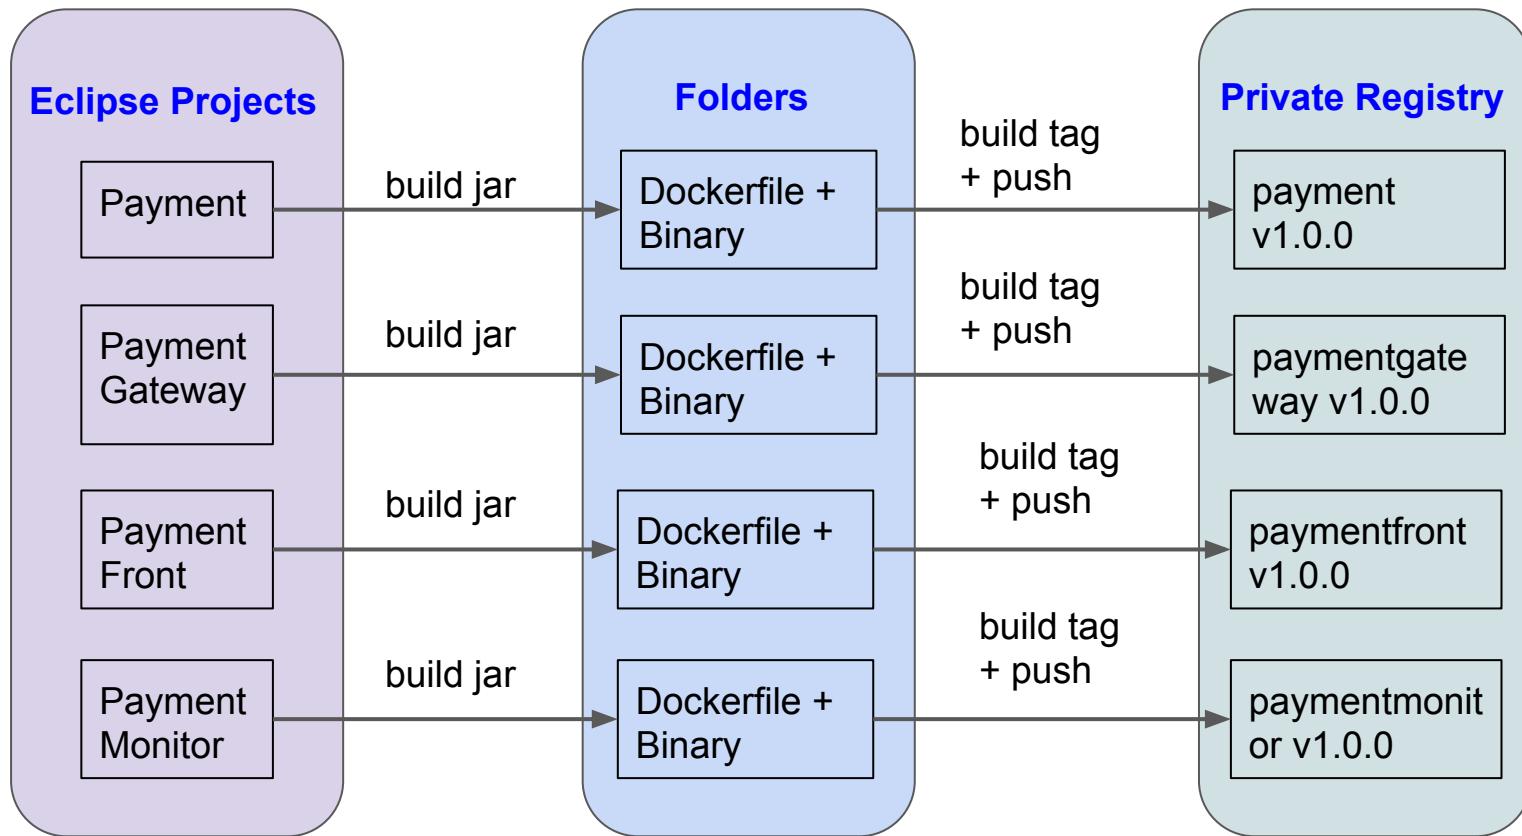
Dockerfile

- A **Dockerfile** is a text based script that **contains instructions and commands** for building the image from the base image.
- Docker reads this Dockerfile when you request a build of an image, executes the instructions, and returns a final image.
- **Docker build** is the process of building a Docker image from a Dockerfile.
- **Docker layer** is each layer in a Docker context represents an instruction included in a Docker image's Dockerfile. The layers can also be referred to as "build steps".
- **Docker build cache**, every time you build a Docker image, each build step is cached. Reuse cached layers that do not change in the image rebuild process to improve the build time.
- Keep track of changes to Dockerfiles with any SCM.

Compare with Object-Oriented



Container Development Example



```
FROM 192.168.99.117:5000/java:8
MAINTAINER AnurakTh
EXPOSE 8900
COPY payment-1.0.0.jar /home/payment.jar
CMD ["/usr/lib/jvm/java-8-openjdk-amd64/bin/java", "-jar", "/home/payment.jar"]
```

Dockerfile Examples

```
FROM 192.168.99.117:5000/java:8
MAINTAINER AnurakTh
EXPOSE 8900
COPY payment-1.0.0.jar /home/payment.jar
CMD ["/usr/lib/jvm/java-8-openjdk-amd64/bin/java", "-jar", "/home/payment.jar"]
```

Java Application

```
1 FROM 10.28.104.174:5000/python:2.7
2 ADD . /code
3 WORKDIR /code
4 RUN chmod +r /code
5 RUN ls /code -la
6 RUN pip install -r requirements.txt
7 CMD python app.py
8
9
```

PHP Application

Docker Image Building

Example commands

- \$ docker build -t myapp .
- \$ docker build -f <Dockerfile name> -t reader .
- \$ docker build .
- \$ docker tag SOURCE_IMAGE[:TAG] TARGET_IMAGE[:TAG]
- \$ docker tag myapp <docker hub account>/myapp:latest
- \$ docker push <docker hub account>/myapp:latest
- *Build with --no-cache and -q*

*** Just version your tags. Every time. ***

*** Do not uses the latest tag in production. ***

*** The ‘latest’ tag does not actually mean latest, it doesn’t mean anything. ***

Dockerfile instructions

- FROM
- RUN
- CMD
- ENTRYPOINT
- LABEL
- EXPOSE
- ENV
- ADD
- COPY
- VOLUME
- USER
- WORKDIR
- ARG
- HEALTHCHECK

Note: The instruction is not case-sensitive.

<https://docs.docker.com/engine/reference/builder/>

FROM Instruction

- The FROM instruction initializes a new build stage and sets the Base Image for subsequent instructions.
- A valid Dockerfile must start with a FROM instruction.
- Can start from a special empty repository named scratch.

```
FROM [--platform=<platform>] <image> [AS <name>]
```

Or

```
FROM [--platform=<platform>] <image>[:<tag>] [AS <name>]
```

Or

```
FROM [--platform=<platform>] <image>[@<digest>] [AS <name>]
```

Run Instruction

- The RUN instruction will **execute any commands** in a new layer on top of the current image and commit the results. The resulting committed image will be used for the next step in the Dockerfile.
- These commands get executed once at build time and get written into your Docker image as a new layer.
- RUN has 2 forms:
 - **RUN <command>** (shell form, the command is run in a shell, which by default is /bin/sh -c on Linux or cmd /S /C on Windows)
 - **RUN ["executable", "param1", "param2"]** (exec form)
- Examples

```
RUN /bin/bash -c 'source $HOME/.bashrc; \
echo $HOME'
```

```
RUN /bin/bash -c 'source $HOME/.bashrc; echo $HOME'
```

```
RUN ["/bin/bash", "-c", "echo hello"]
```

CMD Instruction

- CMD lets you define a default command to run when your container starts.
- The CMD instruction has three forms:
 - `CMD ["executable","param1","param2"]` (exec form, this is the preferred form)
 - `CMD ["param1","param2"]` (as *default parameters to ENTRYPPOINT*)
 - `CMD command param1 param2` (shell form)
- Example
 - the **shell form** of the CMD, then the <command> will execute in /bin/sh -c

```
FROM ubuntu
CMD echo "This is a test." | wc -
```
 - **exec form**, must express the command as a JSON array and give the full path to the executable.

```
FROM ubuntu
CMD ["/usr/bin/wc","--help"]
```
- If you would like your container to run the same executable every time, then you should consider using ENTRYPPOINT in combination with CMD.

ENTRYPOINT Instruction

- An ENTRYPOINT allows you to configure a container that will run as an executable.
- Formats

- The exec form, which is the preferred form:

```
ENTRYPOINT ["executable", "param1", "param2"]
```

- The shell form:

```
ENTRYPOINT command param1 param2
```

- We cannot override the ENTRYPOINT instruction by adding command-line parameters to the docker run command.

Note: There is a way to override the ENTRYPOINT instruction - you need to add the --entrypoint flag prior to the container_name when running the command.

CMD with ENTRYPOINT

- There are some situations in which **combining CMD and ENTRYPOINT** would be the best solution for your Docker container. In such cases, **the executable is defined with ENTRYPOINT, while CMD specifies the default parameter.**
- If you are using both instructions, make sure to keep them in **exec form**.

```
≡ Dockerfile-cmd-entrypoint
 1 FROM ubuntu:18.04
 2 LABEL maintainer=Anurakth
 3 RUN apt-get update
 4 ENTRYPOINT ["echo", "Hello"]
 5 CMD ["World"]
```

LABEL Instruction

- The LABEL instruction adds metadata to an image.
- A LABEL is a key-value pair.
- Format `LABEL <key>=<value> <key>=<value> <key>=<value> ...`
- Examples

```
LABEL "com.example.vendor"="ACME Incorporated"
LABEL com.example.label-with-value="foo"
LABEL version="1.0"
LABEL description="This text illustrates \
that label-values can span multiple lines."
```

```
LABEL multi.label1="value1" multi.label2="value2" other="value3"
```

```
LABEL multi.label1="value1" \
multi.label2="value2" \
other="value3"
```

- To view an image's labels, use the docker image inspect command. You can use the --format option to show just the labels;

```
docker image inspect --format='{{ .Labels }}' myimage
```

```
{
  "com.example.vendor": "ACME Incorporated",
  "com.example.label-with-value": "foo",
  "version": "1.0",
  "description": "This text illustrates that label-values can span multiple lines.",
  "multi.label1": "value1",
  "multi.label2": "value2",
  "other": "value3"
}
```

- **MAINTAINER (deprecated)**

```
LABEL maintainer="SvenDowideit@home.org.au"
```

```
$ docker ps --filter "label=color=blue"
```

EXPOSE Instruction

- The EXPOSE instruction informs Docker that the container **listens on the specified network ports at runtime.**
- Format

```
EXPOSE <port> [<port>/<protocol>...]
```

- Examples

```
EXPOSE 80/tcp  
EXPOSE 80/udp
```

- To actually publish the port when running the container, use the -p flag on docker run

```
docker run -p 80:80/tcp -p 80:80/udp ...
```

ENV Instruction

- The ENV instruction sets the environment variable <key> to the value <value>.
- Format

```
ENV <key>=<value> ...
```

- Examples

```
ENV MY_NAME="John Doe"  
ENV MY_DOG=Rex\ The\ Dog  
ENV MY_CAT=fluffy
```

```
ENV MY_NAME="John Doe" MY_DOG=Rex\ The\ Dog \  
MY_CAT=fluffy
```

```
ENV MY_VAR my-value
```

The alternative syntax is supported
for backward compatibility.

- The environment variables set using ENV will persist when a container is run from the resulting image.
- You can view the values using docker inspect, and change them using
 - docker run --env <key>=<value>

ADD Instruction

- The ADD instruction copies new **files, directories or remote file URLs** from <src> and adds them to the filesystem of the image at the path <dest>.
- Format

```
ADD [--chown=<user>:<group>] <src>... <dest>
ADD [--chown=<user>:<group>] [<src>, ... <dest>]
```

- Example

```
ADD test.txt relativeDir/
```

```
ADD test.txt /absoluteDir/
```

```
ADD hom* /mydir/
```

```
ADD hom?.txt /mydir/
```

```
ADD --chown=55:mygroup files* /somedir/
ADD --chown=bin files* /somedir/
ADD --chown=1 files* /somedir/
ADD --chown=10:11 files* /somedir/
```

- If <src> is a local tar archive in a recognized compression format (identity, gzip, bzip2 or xz) then it is unpacked as a directory.

```
Users > ktb_user > Documents > dockerTmp > JbossApp1 > Dockerfile > FROM
1  FROM 192.168.99.117:5000/jboss/wildfly:v1.0.1
2  ADD node-info.war /opt/jboss/wildfly/standalone/deployments/
3  ADD node-info.war.dodeploy /opt/jboss/wildfly/standalone/deployments/
```

Copy Instruction

- The COPY instruction copies new **files or directories** from <src> and adds them to the filesystem of the container at the path <dest>.
- Formats

```
COPY [--chown=<user>:<group>] <src>... <dest>
COPY [--chown=<user>:<group>] [<src>, ... <dest>]
```

- Examples

```
COPY hom* /mydir/    COPY hom?.txt /mydir/
```

```
COPY --chown=55:mygroup files* /somedir/
COPY --chown=bin files* /somedir/
COPY --chown=1 files* /somedir/
COPY --chown=10:11 files* /somedir/
```

VOLUME Instruction

- The VOLUME instruction creates a mount point with the specified name and marks it as holding externally mounted volumes from native host or other containers.
- Formats

```
VOLUME ["/data"]
```

- Examples

```
FROM ubuntu
RUN mkdir /myvol
RUN echo "hello world" > /myvol/greeting
VOLUME /myvol
```

- Override volume name during docker run with -v option
 - docker run -v mydata:/myvol <image-name>

USER Instruction

- The USER instruction sets the **user name (or UID)** and optionally the **user group (or GID)** to use when running the image and for any RUN, CMD and ENTRYPOINT instructions that follow it in the Dockerfile.
- Formats

```
USER <user>[:<group>]
```

```
USER <UID>[:<GID>]
```

- Examples

```
FROM ubuntu:latest
RUN useradd -r -u 1001 -g appuser appuser
USER appuser
ENTRYPOINT ["sleep", "infinity"]
```

- Another option is to run a docker container and specify the username or uid, and also the group name or gid at runtime.
 - *\$ docker run -d --user 1001 ubuntu:latest sleep infinity*

WORKDIR Instruction

- The WORKDIR instruction sets the working directory for any RUN, CMD, ENTRYPOINT, COPY and ADD instructions that follow it in the Dockerfile.
- If the WORKDIR doesn't exist, it will be created even if it's not used in any subsequent Dockerfile instruction.
- Formats

```
WORKDIR /path/to/workdir
```

- Examples

```
WORKDIR /a  
WORKDIR b  
WORKDIR c  
RUN pwd
```

The output of the final `pwd` command in this Dockerfile would be `/a/b/c`.

ARG

- ARG is the only instruction that may precede FROM in the Dockerfile
- The ARG instruction defines a variable that users can **pass at build-time** to the builder with the docker build command using the **--build-arg <varname>=<value>**
- Formats

```
ARG <name>[=<default value>]
```

Predefined ARGs

- Examples

```
FROM ubuntu
ARG CONT_IMG_VER
ENV CONT_IMG_VER=v1.0.0
RUN echo $CONT_IMG_VER
```

```
FROM busybox
ARG user1=someuser
ARG buildno=1
# ...
```

```
$ docker build --build-arg CONT_IMG_VER=v2.0.1 .
```

- HTTP_PROXY
- http_proxy
- HTTPS_PROXY
- https_proxy
- FTP_PROXY
- ftp_proxy
- NO_PROXY
- no_proxy

--build-arg HTTP_PROXY=http://user:pass@proxy.lon.example.com

HEALTHCHECK

- The HEALTHCHECK instruction tells Docker how to test a container to check that it is still working.
- When a container has a healthcheck specified, it has a health status in addition to its normal status.
- Formats

```
HEALTHCHECK [OPTIONS] CMD command (check container health by running a command inside the container)  
HEALTHCHECK NONE (disable any healthcheck inherited from the base image)
```

The options that can appear before `CMD` are:

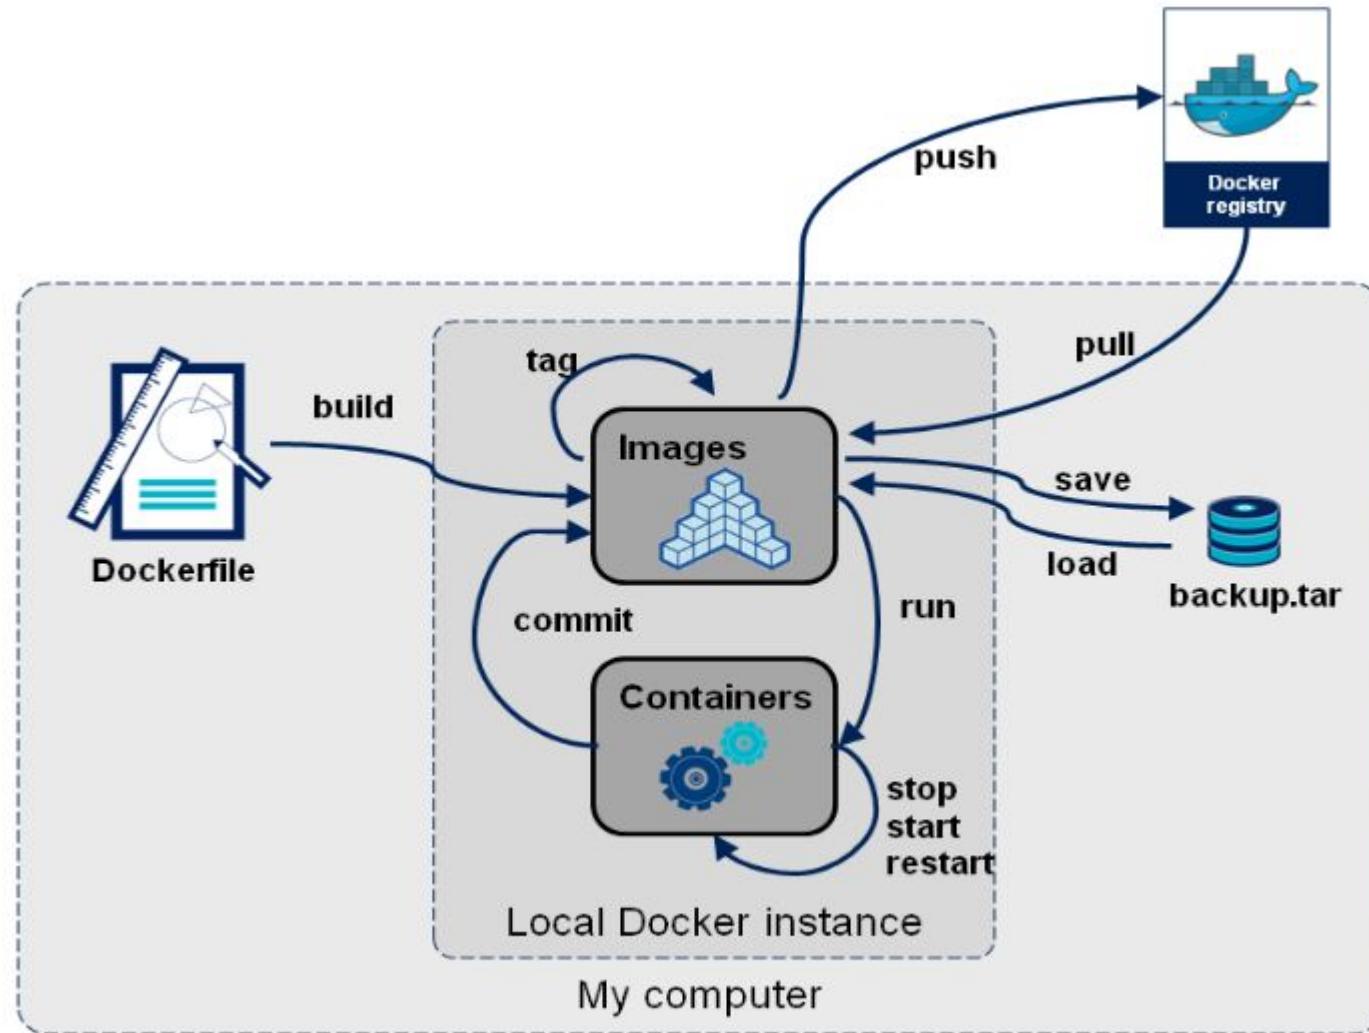
- `--interval=DURATION` (default: `30s`)
- `--timeout=DURATION` (default: `30s`)
- `--start-period=DURATION` (default: `0s`)
- `--retries=N` (default: `3`)

- Examples

```
HEALTHCHECK --interval=5m --timeout=3s \  
CMD curl -f http://localhost/ || exit 1
```

- 0: success - the container is healthy and ready for use
- 1: unhealthy - the container is not working correctly
- 2: reserved - do not use this exit code

Container Development



Java Dockerfile

```
# select parent image
FROM maven:3.6.3-jdk-8

# copy the source tree and the pom.xml to our new container
COPY ./ ./
```



```
# package our application code
RUN mvn clean package
```



```
# set the startup command to execute the jar
CMD ["java", "-jar", "target/demo-0.0.1-SNAPSHOT.jar"]
```

Go Dockerfile

```
FROM golang:latest
```

```
WORKDIR /myap
```

```
COPY welcome.go .
```

```
RUN go build -o welcome .
```

```
ENTRYPOINT ["./welcome"]
```

Docker Build Cache

Show docker disk usage

```
$ docker system df
```

Remove Docker Build Cache

```
$ docker images -a
```

```
$ docker builder prune
```

Dangling Images

```
$ docker images --filter "dangling=true"
```

```
$ docker rmi $(docker images -q --filter "dangling=true")
```

```
$ docker system prune
```

```
$ docker system prune --help
```

Multi-stage builds

- Introduces in the 17.05 version.
- Very useful to reduce the image size of the production image.
- A multistage Dockerfile is a Dockerfile that has multiple FROM instructions.
- Each FROM instruction marks a new build stage whose final layer may be referenced in a downstream stage.
- **FROM BASE_IMAGE <AS STAGE_NAME>**
- **COPY --from=STAGE_NAME**

Multi-stage Java Example

```
# the first stage of our build will use a maven 3.6.1 parent image
FROM maven:3.6.1-jdk-8-alpine AS MAVEN_BUILD
```

```
# copy the pom and src code to the container
COPY ./ ./
```

```
# package our application code
RUN mvn clean package
```

```
# the second stage of our build will use open jdk 8 on alpine 3.9
FROM openjdk:8-jre-alpine3.9
```

```
# copy only the artifacts we need from the first stage and discard the rest
COPY --from=MAVEN_BUILD /target/demo-0.0.1-SNAPSHOT.jar /demo.jar
```

```
# set the startup command to execute the jar
CMD ["java", "-jar", "/demo.jar"]
```

Multi-stage Go Example

```
# Start from latest golang parent image
FROM golang:latest AS builder
# Set the working directory
WORKDIR /myapp
# Copy source file from current directory to container
COPY welcome.go .
# Build the application
RUN go build -o welcome .

# Start from latest alpine parent image
FROM alpine:latest AS runtime
# Set the working directory
WORKDIR /myap
# Copy helloworld app from current directory to container
COPY --from=builder /myapp/welcome .
# Run the application
ENTRYPOINT ["../welcome"]
```

Dockerfile best practices

<https://docs.docker.com/develop/dev-best-practices/>

https://docs.docker.com/develop/develop-images/dockerfile_best-practices/

Docker Compose

Docker Compose Overview

- Docker Compose is a **tool for defining and running multi-container Docker applications.**
- Use a **Compose file** to configure your application's **services**.
- Using a single command, you create and start all the services from your configuration.
- Using Docker Compose is basically a three-step process.
 1. Define your app's environment with a **Dockerfile** so it can be reproduced anywhere.
 2. Define the services that make up your app in **docker-compose.yml** so they can be run together in an isolated environment.
 3. Lastly, run **docker-compose up** and Compose will start and run your entire app.
- Docker for Mac, Docker for Windows, and **Docker Toolbox** include Docker Compose.

Compose and Docker compatibility matrix

There are several versions of the Compose file format – 1, 2, 2.x, and 3.x. The table below is a quick look. For full details on what each version includes and how to upgrade, see [About versions and upgrading](#).

This table shows which Compose file versions support specific Docker releases.

Compose file format	Docker Engine release
3.8	19.03.0+
3.7	18.06.0+
3.6	18.02.0+
3.5	17.12.0+
3.4	17.09.0+
3.3	17.06.0+
3.2	17.04.0+
3.1	1.13.1+
3.0	1.13.0+
2.4	17.12.0+
2.3	17.06.0+

Try It!

Try Docker Compose

```
version: "2"
services:
  reader:
    image: anutech2001/reader
    ports:
      - "8010:8010"
    links:
      - bookstore
      - zipkin
    labels:
      kompose.service.type: NodePort
  bookstore:
    image: anutech2001/bookstore
    ports:
      - "8011:8011"
    labels:
      kompose.service.type: NodePort
    links:
      - zipkin
```

```
zipkin:
  image: openzipkin/zipkin
  ports:
    - "9411:9411"
  labels:
    kompose.service.type: NodePort
```

Play with docker compose

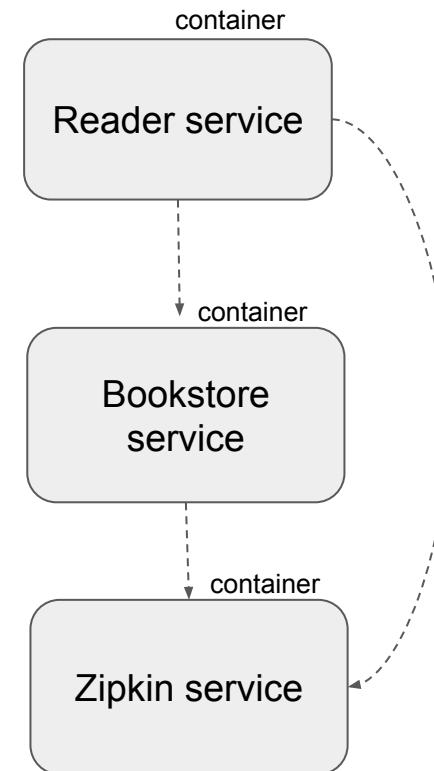
- docker-compose up -d(optional)
- docker-compose logs -f
- docker-compose scale reader=2
- docker-compose stop
- Docker-compose down

<https://docs.docker.com/compose/reference/>

Try Docker Compose (con't)

```
version: "2"
services:
  reader:
    image: anutech2001/reader
    ports:
      - "8010:8010"
    links:
      - bookstore
    ....
```

docker-compose up



Docker Swarm

What is a Swarm?

- The **cluster management and orchestration** features embedded in the Docker Engine
- **A swarm is a cluster of Docker engines**, or **nodes**, where you deploy services. The Docker Engine CLI and API include commands to manage swarm nodes (e.g., add or remove nodes), and deploy and orchestrate services across the swarm.
- A **node** is an instance of the Docker engine participating in the swarm. You can also think of this as a Docker node.
- Production swarm deployments typically include Docker nodes distributed across multiple physical and cloud machines.

Swarm Mode Features

- **Cluster management integrated with Docker Engine**

Don't need additional orchestration software to create or manage a swarm.

- **Scaling**

swarm manager automatically adapts by adding or removing tasks to maintain the desired state

- **Desired state reconciliation**

The swarm manager node constantly monitors the cluster state and reconciles any differences between the actual state and your expressed desired state.

- **Multi-host networking**

You can specify an overlay network for your services. The swarm manager automatically assigns addresses to the containers on the overlay network when it initializes or updates the application.

- **Service discovery**

Swarm manager nodes assign each service in the swarm a unique DNS name and load balances running containers.

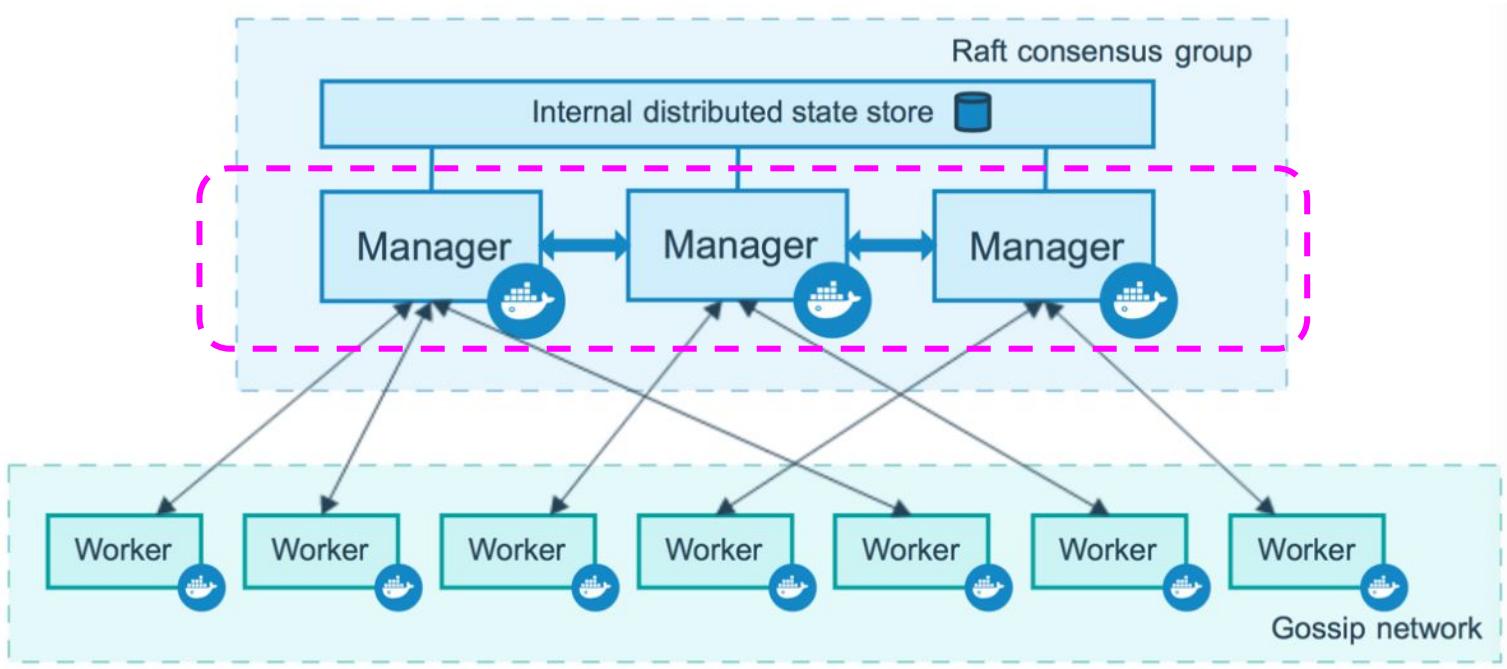
- **Load balancing**

- **Secure by default**

- **Rolling updates**

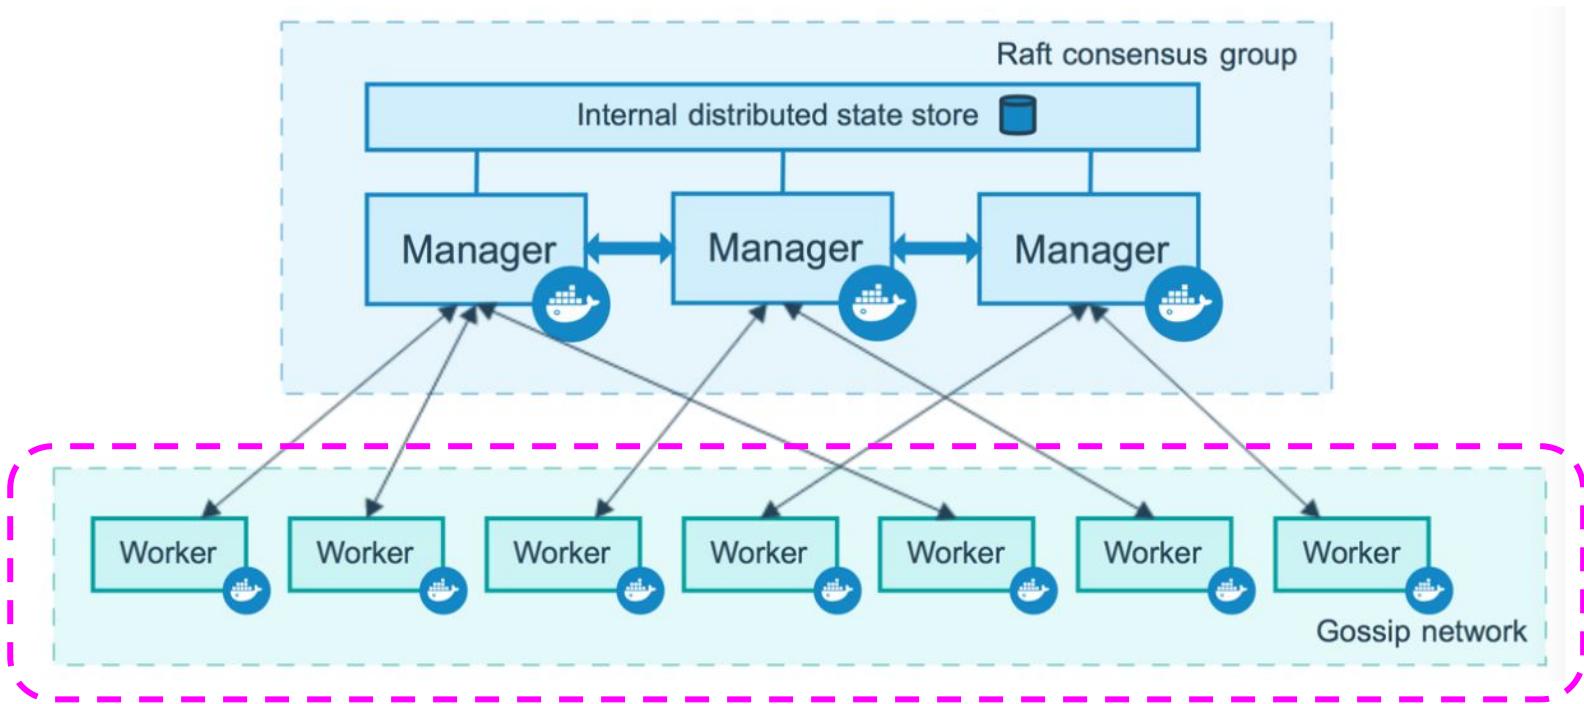
A Manager Node

- Manager nodes perform the **orchestration and cluster management** functions required to maintain the desired state of the swarm.
- To deploy your application to a swarm, you submit a service definition to a manager node.
- The manager node dispatches units of work called tasks to **worker nodes**.
- As worker nodes or just a manager-only nodes.



A Worker Node

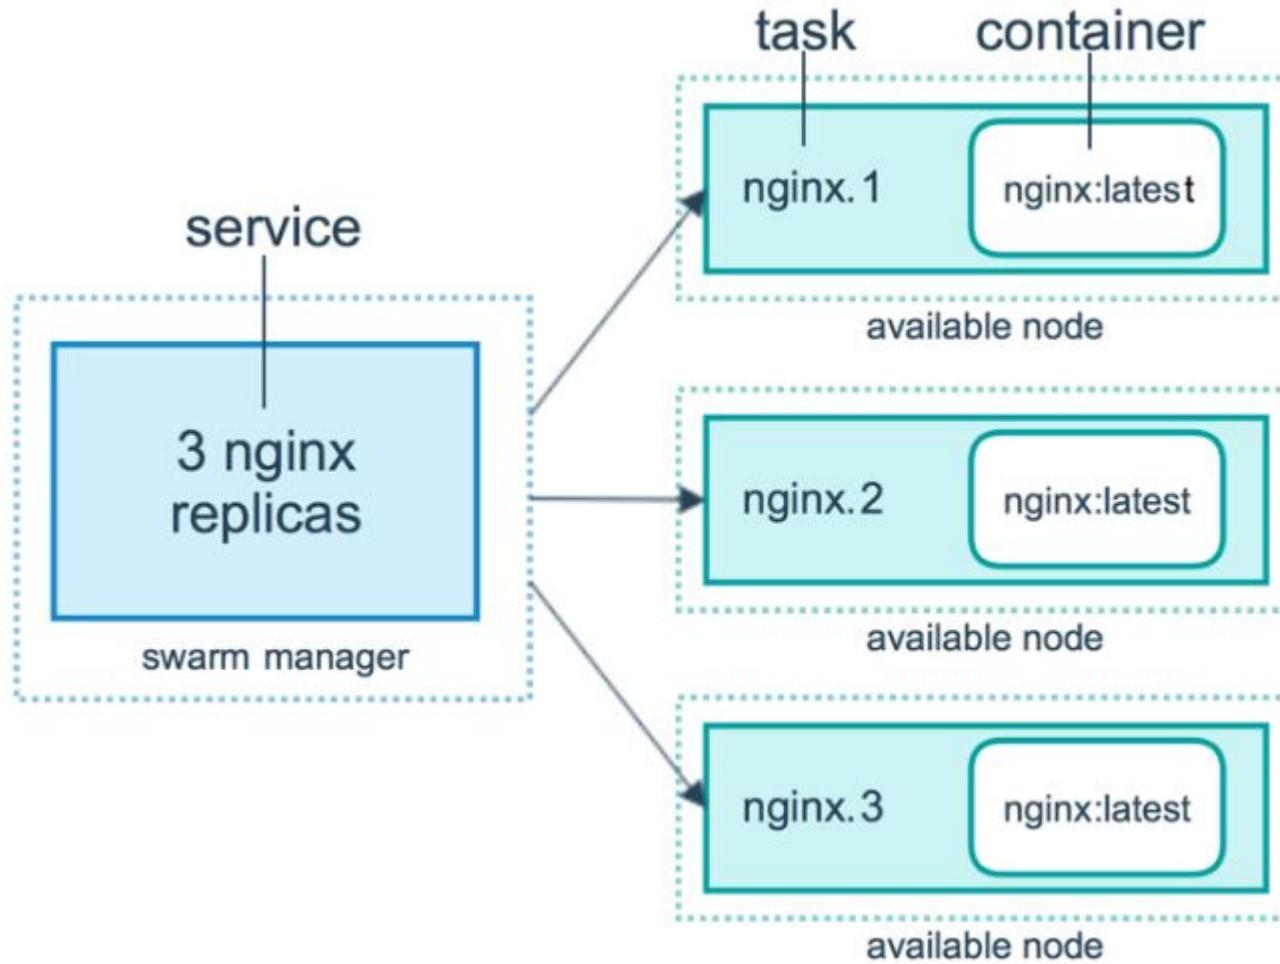
- **Worker nodes** receive and execute tasks dispatched from manager nodes.
- By default manager nodes also run services as worker nodes (can turn it off).
- You can promote a worker node to be a manager.



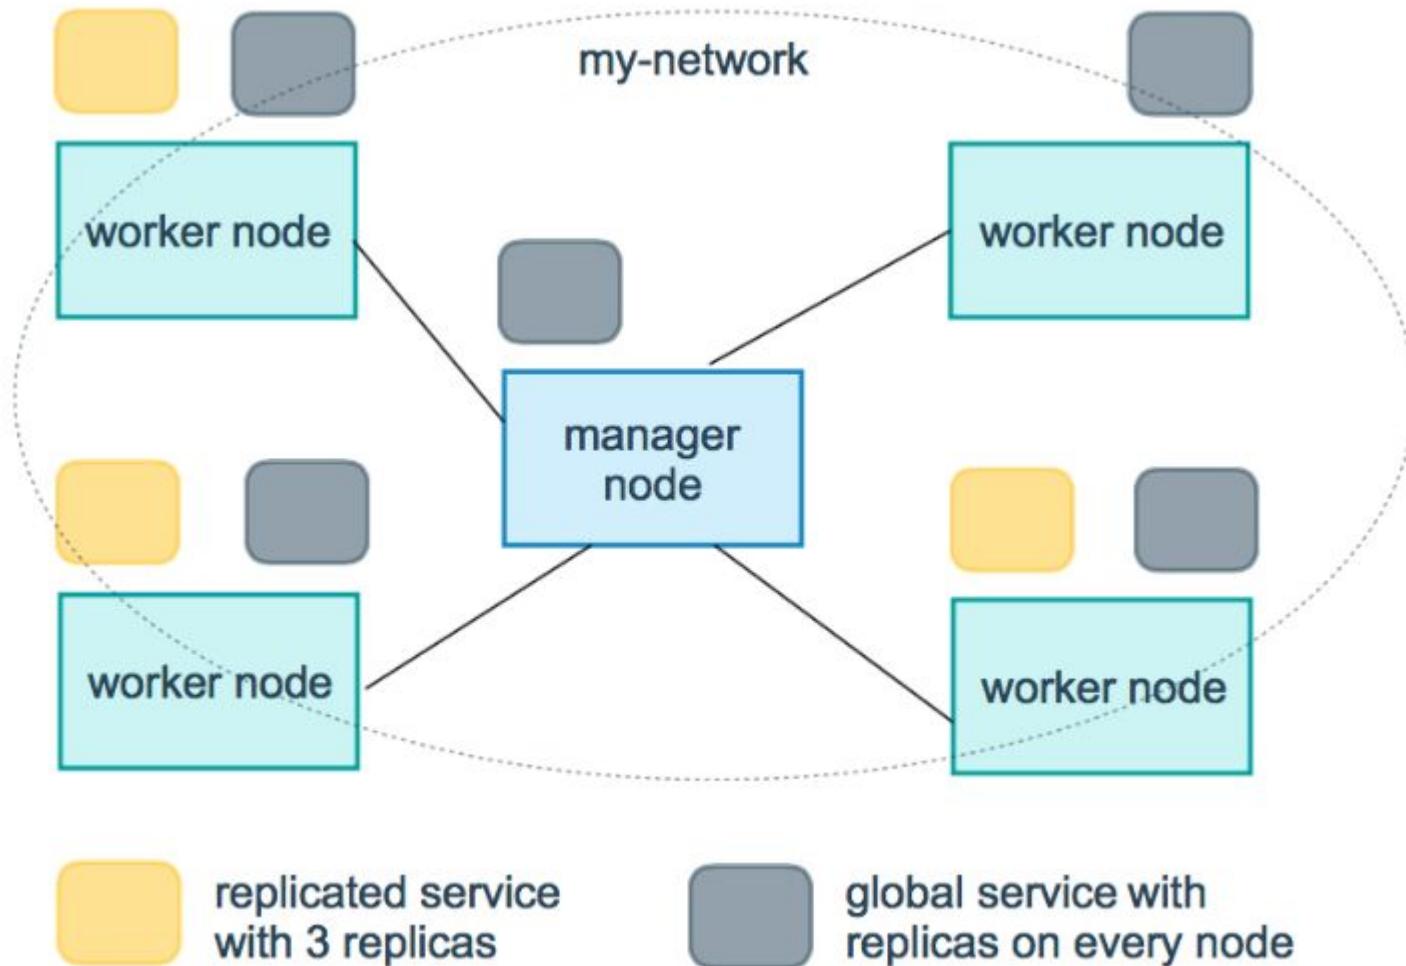
Services and tasks

- A service is the definition of the tasks to execute on the worker nodes.
- **Changes from “*docker run*” to “*docker create service*” command**
- When you create a service, you specify which container image to use and which commands to execute inside running containers.
- **Replicated services model**
 - the swarm manager distributes a specific number of replica tasks among the nodes based upon the scale you set in the desired state.
- **Global services model**
 - the swarm runs one **task** for the service on every available node in the cluster.
- **Task** is the atomic scheduling unit of swarm.
- Manager nodes assign tasks to worker nodes according to the number of replicas set in the service scale.

Services, tasks and Containers



Replicated services and Global services



Load balancing

- The swarm manager uses **ingress load balancing** to expose the services
- The swarm manager can automatically assign the service a PublishedPort or you can configure a PublishedPort for the service.
- the swarm manager assigns the service a port in the 30000-32767 range.
- External components can access the service on the PublishedPort of any node in the cluster **whether or not the node is currently running the task for the service.**
- Swarm mode has an **internal DNS component** that automatically assigns each service in the swarm a DNS entry
- The swarm manager uses **internal load balancing** to distribute requests among services within the cluster based upon the DNS name of the service.

Play with Swarm Mode

- **Create a swarm**
- **Join nodes to a swarm** - as a worker and a manager node
- **Manage nodes in a swarm** - list, inspect, update a node and leave the swarm
- **Deploy services to a swarm** - create, update, config, control and scale a service
- <https://labs.play-with-docker.com/>

Try It!

Create a swarm

- Create swarm (@Manager)

```
$ docker swarm init  
Swarm initialized: current node (dxn1zf6l6lqsbljosjja83ngz) is now a manager.
```

To add a worker to this swarm, run the following command:

```
docker swarm join \  
--token SwMTKN-1-49nj1cmql0jkz5s954yi3oex3nedyz0fb0xx14ie39trti4wxv-8vxv8rssmk743ojnwacrr2e7c \  
192.168.99.100:2377
```

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

```
$ docker swarm init
```

```
$ docker swarm init --advertise-addr <MANAGER-IP> -- กรณีมีหลาย IP
```

Try It!

Join nodes to a swarm - as worker node

- Join as a worker node (@Node)

```
$ docker swarm join \
--token SWMTKN-1-49nj1cmql0jkz5s954yi3oex3nedyz0fb0xx14ie39trti4wxv-8vxv8rssmk743ojnwacrr2e7c \
192.168.99.100:2377

This node joined a swarm as a worker.
```

- Check join token for worker (@Manager)

```
$ docker swarm join-token worker

To add a worker to this swarm, run the following command:

docker swarm join \
--token SWMTKN-1-49nj1cmql0jkz5s954yi3oex3nedyz0fb0xx14ie39trti4wxv-8vxv8rssmk743ojnwacrr2e7c \
192.168.99.100:2377
```

Try It!

Join nodes to a swarm - as manager node

- Check join token for manager (@Manager)

```
$ docker swarm join-token manager
```

To add a manager to this swarm, run the following command:

```
docker swarm join \  
--token SWMTKN-1-61ztec5kyafptydic6jfc1i33t37f1c14nuipzcusor96k7kby-5vy9t8u35tuqm7vh671rz9xp6 \  
192.168.99.100:2377
```

- Join as a manager node (@Node)

```
$ docker swarm join \  
--token SWMTKN-1-61ztec5kyafptydic6jfc1i33t37f1c14nuipzcusor96k7kby-5vy9t8u35tuqm7vh671rz9xp6 \  
192.168.99.100:2377
```

This node joined a swarm as a manager.

Try It!

Manage nodes in a swarm

- Create 1 manager and 2 worker nodes.
- List nodes
 - \$ docker node ls (@Manager)
- Inspect an individual node
 - \$ docker node inspect <NODE> (@Manager)
- Change node availability
 - \$ docker node update --availability drain <NODE> (@Manager)
 - Set DRAIN availability prevents a node from receiving new tasks from the swarm manager.
 - \$ docker node update --availability active <NODE> (@Manager)
 - Set ACTIVE availability to return the drained node to an active state
- Add or remove label metadata
 - \$ docker node update --label-add <key>=<value>... <NODE> (@Manager)
- Promote or demote a node
 - \$ docker node promote <NODE>... (@Manager - promote a node to be manager)
 - \$ docker node demote <NODE>... (@Manager - demote a node from a manager node)
- Leave the swarm
 - \$ docker swarm leave (@Node)

Deploy services to a swarm

- Create a service
 - \$ docker service create <IMAGE> (@Manager)
- To list the service
 - \$ docker service ls (@Manager)
- Add an overlay network, use overlay networks to connect one or more services within the swarm.
 - \$ docker network create --driver overlay my-network (@Manager)
- Configure update behavior
 - <https://docs.docker.com/engine/swarm/services/#configure-update-behavior>

Try It!

Example - initial swarm

- Create 3 docker machines name manager, service1, db1
- docker swarm init --advertise-addr <manager ip> (@Manager)
- docker network create --driver overlay countnet (@Manager)
- OR docker network create --driver overlay --subnet=10.0.9.0/24 countnet (@Manager)
- docker swarm join --token <token> <manager ip>:2377 (@Node)
- docker node update --label-add.nodeType=service service1(@Manager)
- docker node update --label-add.nodeType=db db1(@Manager)
- docker node update --availability drain manager(@Manager)
- docker node ls
- docker node update --availability active manager(@Manager)
- docker node ls

Try It!

Example - Manage Service

- docker service create --name myweb --replicas 3 --publish published=8080,target=80 nginx
- docker service ls
- docker service ps myweb
- docker service inspect --pretty myweb
- docker service logs -f myweb
- curl <http://localhost:8080> -- *try it many times.*
- See the logs. Who gets the requests?
- docker stop <one of myweb container>
- docker service ps myweb -- *What happens?*
- Check ingress network
- docker service rm myweb
- docker service update --image nginx:<version> nginx

Try It!

Example - Manage Service(con't)

- docker network create --driver overlay countnet
- docker network create --driver overlay --subnet=10.0.9.0/24 countnet
- docker node update --label-add.nodeType=service service1
- docker node update --label-add.nodeType=db db1
- docker service create --name redis --constraint node.labels.nodeType==db --network countnet 192.168.99.117:5000/redis
- docker service create --name countapp --constraint node.labels.nodeType==service -p 5001:5000 --network countnet --mount type=bind,source=/c/Users/<username>/code/,target=/code anutech2001/countapp:v1.0.2
- Browse to <http://<node ip>:5001>

Note: use [-P](#) option in docker run command to tell docker to publish all exposed ports to random ports.

Try It!

Example - Web Application with Jboss

- docker service create --name jbossapp1 -p 8800:8080 -p 9990:9990 anutech2001/jbossapp1:v1.0.0
- docker service scale jbossapp1=5
- Docker service ls
- docker service ps jbossapp1
- Browse to `http://<manager ip>:8800/node-info` then observe the value of ***Hostname*** when you try to refresh the page every 5-7 seconds.

```
Hostname: 2cfaa8f13ab9
Java Runtime: OpenJDK Runtime Environment 1.8.0_101-b13
OS: Linux amd64 4.4.27-boot2docker
```

- docker service rm jbossapp1
- docker service create --name jbossapp1 --constraint node.labels.nodeType==service -p 8800:8080 -p 9990:9990 anutech2001/jbossapp1:v1.0.0

Note: user for login to jboss admin is “**admin**” and password is “**Admin#70365**”.

Docker Stack

A Stack

- A stack is a group of interrelated services that share dependencies, and can be orchestrated and scaled together.
- ***docker stack*** is a command that's embedded into the Docker CLI.
- It lets you manage a cluster of Docker containers through **Docker Swarm**.
- Docker and docker service command still can use within stack service.

Stack Commands

Command	Description
<code>docker stack deploy</code>	Deploy a new stack or update an existing stack
<code>docker stack ls</code>	List stacks
<code>docker stack ps</code>	List the tasks in the stack
<code>docker stack rm</code>	Remove one or more stacks
<code>docker stack services</code>	List the services in the stack

docker stack deploy

Options

Name, shorthand	Default	Description
--compose-file , -c		API 1.25+ Path to a Compose file, or “-” to read from stdin
--namespace		Kubernetes Kubernetes namespace to use
--prune		API 1.27+ Swarm Prune services that are no longer referenced
--resolve-image	always	API 1.30+ Swarm Query the registry to resolve image digest and supported platforms (“always” “changed” “never”)
--with-registry-auth		Swarm Send registry authentication details to Swarm agents
--kubeconfig		Kubernetes Kubernetes config file
--orchestrator		Orchestrator to use (swarm kubernetes all)

Voting app

```
version: "3.8"
services:

  redis:
    image: redis:alpine
    ports:
      - "6379"
    networks:
      - frontend
    deploy:
      replicas: 2
      update_config:
        parallelism: 2
        delay: 10s
      restart_policy:
        condition: on-failure

  db:
    image: postgres:9.4
    environment:
      POSTGRES_HOST_AUTH_METHOD: 'trust'
      POSTGRES_PASSWORD: 'password'
    volumes:
      - db-data:/var/lib/postgresql/data
    networks:
      - backend
    deploy:
      placement:
        max_replicas_per_node: 1
      constraints:
        - "node.role==manager"

  vote:
    image:
      dockersamples/examplevotingapp_vote:before
    ports:
      - "5000:80"
    networks:
      - frontend
    depends_on:
      - redis
    deploy:
      replicas: 2
      update_config:
        parallelism: 2
      restart_policy:
        condition: on-failure

  result:
    image:
      dockersamples/examplevotingapp_result:before
    ports:
      - "5001:80"
    networks:
      - backend
    depends_on:
      - db
    deploy:
      replicas: 1
      update_config:
        parallelism: 2
        delay: 10s
      restart_policy:
        condition: on-failure

  worker:
    image: dockersamples/examplevotingapp_worker
    networks:
      - frontend
      - backend
    deploy:
      mode: replicated
      replicas: 1
      labels: [APP=VOTING]
      restart_policy:
        condition: on-failure
        delay: 10s
        max_attempts: 3
        window: 120s
      placement:
        constraints:
          - "node.role==manager"

  visualizer:
    image: dockersamples/visualizer:stable
    ports:
      - "8080:8080"
    stop_grace_period: 1m30s
    volumes:
      - "/var/run/docker.sock:/var/run/docker.sock"
    deploy:
      placement:
        constraints:
          - "node.role==manager"

  networks:
    frontend:
    backend:

  volumes:
    db-data:
```

Docker Stack Demo

Deploy on Swarm

<https://labs.play-with-docker.com/>

```
$ docker stack deploy -c docker-compose.yml app
```

Deploy on Kubernetes

```
$ kubectl create namespace training
```

```
$ docker stack deploy --orchestrator kubernetes --namespace training -c docker-compose.yml app
```

<https://docs.docker.com/docker-for-mac/kubernetes/>

Services, Tasks and Containers

```
[manager1] (local) root@192.168.0.7 ~
$ docker stack services app
ID          NAME      MODE      REPLICAS      IMAGE
jgx183grha6v  app_db    replicated  1/1 (max 1 per node)  postgres:9.4
lyfrhh7vk2ak  app_redis  replicated  2/2
rjxaedzduqu27  app_result  replicated  1/1
noin3j6rnmxo9  app_visualizer  replicated  1/1
td5igmf5zwlm  app_vote    replicated  2/2
vshrl1dqzzcg3  app_worker   replicated  1/1
[manager1] (local) root@192.168.0.7 ~
$ docker stack ps app
ID          NAME      IMAGE      NODE      DESIRED STATE     CURRENT STATE      ERROR      PORTS
mjoda1fq0y01  app_db.1  postgres:9.4  manager1  Running       Running  23 minutes ago
pa3hzeub3a21  app_redis.1  redis:alpine  worker1    Running       Running  27 minutes ago
ijucy9c4399d  app_redis.2  redis:alpine  worker1    Running       Running  27 minutes ago
78ftpjaxaken  app_result.1  dockersamples/exempllevotingapp_result:before  worker1    Running       Running  27 minutes ago
t7whxn5po3ds  app_visualizer.1  dockersamples/visualizer:stable  manager1  Running       Running  23 minutes ago
he1nftg9ygeo  app_vote.1  dockersamples/exempllevotingapp_vote:before  worker1    Running       Running  27 minutes ago
hkhw9xj0ef19  app_vote.2  dockersamples/exempllevotingapp_vote:before  worker1    Running       Running  27 minutes ago
4erdsxcdlyr7  app_worker.1  dockersamples/exempllevotingapp_worker:latest  manager1  Running       Running  22 minutes ago
[manager1] (local) root@192.168.0.7 ~
$ docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS                 NAMES
f1c7fe5b1e03        dockersamples/exempllevotingapp_worker:latest  "/bin/sh -c 'dotnet ..."  25 minutes ago      Up 25 minutes           app_worker.1.4erd
sxcldlyr75qpzhkky1pkvf
a8f6dd385acb        dockersamples/visualizer:stable            "npm start"            25 minutes ago      Up 25 minutes           8080/tcp      app_visualizer.1.
t7whxn5po3ds54sfhluy4hw
c7270b8f6fc3        postgres:9.4                         "docker-entrypoint.s..."  25 minutes ago      Up 25 minutes           5432/tcp      app_db.1.mjodalfq
0y01bh5eysg3m436f
```

```
$ docker inspect f1c
```

```
"Labels": {
  "com.docker.stack.namespace": "app",
  "com.docker.swarm.node.id": "6ndizxgogzjuisrfzjhcu43aa",
  "com.docker.swarm.service.id": "vshrl1dqzzcg301arpntj9u2kx",
  "com.docker.swarm.service.name": "app_worker",
  "com.docker.swarm.task": "",
  "com.docker.swarm.task.id": "4erdsxcdlyr75qpzhkky1pkvf",
  "com.docker.swarm.task.name": "app_worker.1.4erdsxcdlyr75qpzhkky1pkvf"
```

Docker Desktop Kubernetes



The screenshot shows the 'Kubernetes' tab of the Docker Desktop settings in a modern UI style. The title bar says 'Kubernetes v1.19.3'. On the left is a sidebar with the following items:

- General
- Resources
- Docker Engine
- Experimental Features
- Kubernetes (highlighted)

The main area has a checked checkbox for 'Enable Kubernetes' with the subtext: 'Start a Kubernetes single-node cluster when starting Docker Desktop.' Below it is an unchecked checkbox for 'Deploy Docker Stacks to Kubernetes by default' with the subtext: 'Make Kubernetes the default orchestrator for "docker stack" commands'. There is also an unchecked checkbox for 'Show system containers (advanced)' with the subtext: 'Show Kubernetes internal containers when using Docker commands.'

A red-bordered button labeled 'Reset Kubernetes Cluster' is at the bottom. A note below it says: 'All stacks and Kubernetes resources will be deleted.'

Bye.